



The Complete Developers Guide

Section 1

INSTALLING & EXPLORING NODE.JS

Installing Node.js

We can visit <https://nodejs.org/en/> homepage to download node.js for our operating system. There are two versions:

The LTS version (Long Term Support) and the Current release version.

It is advisable to download the Current release, as this will contain all the latest features of Node.js. As at August 2019, the latest version is 12.7.0 Current.

We can test by running a simple terminal command on our machine to see if node.js was actually installed correctly:

```
:~$ node -v
```

This will return a version of node that is running on your machine and will indicate that node was installed.

If we receive command not found, then this means node was not installed and we need to reinstall it.

Installing A Code Editor

A text editor is required so that we can start writing out our node.js scripts. There are a couple of open source code editors (*i.e. free*) that we can use such as Atom, Sublime, notepad++ and Visual Studio Code to name a few.

It is highly recommended to download and install Visual Studio Code as your code editor of choice as it contains many great features, themes and extensions to really customise the editor to fit your needs. We can also debug our node.js scripts inside of the editor which is also a great feature.

Links:

<https://code.visualstudio.com/>

<https://atom.io>

<https://www.sublimetext.com/>

What is Node.js?

Node.js came about when the original developers took JavaScript (*something that ran only on browsers*) and they allowed it to run as a stand alone process on our machine. This means that we can use the JavaScript programming language outside of the browsers and build things such as a web server which can access the file system and connect to databases. Therefore, JavaScript developers could now use

JavaScript as a server side language to create web servers, command line interfaces (CLI), application backends and more. Therefore, at a high level, node.js is a way to run JavaScript code on the server as opposed to being forced to run on only the client.

"Node.js® is a JavaScript runtime built on Chrome's V8 JavaScript engine."

There are all sorts of JavaScript engines out there running all major browsers, however, the Chrome V8 engine is the same engine that runs the Google Chrome browser. The V8 engine is a open source project. The job of the engine is to take JavaScript code and compile it down into machine code that the machine can actually execute. The V8 engine is written in the C++ language.

Node.js is not a programming language. The runtime is something that provides custom functionality i.e. various tools and libraries specific to an environment. So in the case of Chrome, the runtime provides V8 with various objects and functions that allow JavaScript developers in the chrome browser to do things like add a button click events or manipulate the DOM. The node runtime provides various tools that node developers need such as libraries for setting up web-servers integrating with the file system.

The JavaScript is provided to the V8 engine as part of the runtime. JavaScript does not know the methods e.g. read file from disk or get item from local storage but C++ does know the methods. So we have a series of methods that can be used in our JavaScript code which are, in reality, just running C++ code behind the scenes.

Therefore, we have C++ bindings to V8 which allows JavaScript, in essence, to do anything that C++ can do. Below is a visualisation to demonstrate exactly what is happening with the V8 JavaScript Engine:

JavaScript (Chrome)	C++
localStorage.getItem	Some C++ function
document.querySelector	Some C++ function

JavaScript (Node.js)	C++
fs.readFile	Some C++ function
os.platform	Some C++ function



We can run the following code in our terminal

```
:~$ node
```

What we get by running this command is a little place where we can run individual node JavaScript statements which is also known as *repl* (*read eval print loop*). These are not bash commands. All of the core JavaScript features we are used to are still available when using node.js because those are provided by the V8 engine itself.

It is important to note that some of the JavaScript features available in the browser are not available in node. For example, Chrome has a object called window which provides all the different methods and properties we have access to. This makes sense in the context of JavaScript in the browser because we actually have a window to work with. This will not work on node because window is something specifically provided by the Chrome runtime when JavaScript is running in the application. Node does not have a window and it does not need window and therefore window is not provided.

With node we have something similar to window; we have a variable called global:

```
:~$ node  
[> global
```

Global stores a lot of the global things we can access such as the methods and properties available to us. The browser does not have access to the global variable.

Another difference between the browser runtime and the node runtime, is that the browser has access to something called a document. The document allows us to work with the DOM (document object manager) which allows us to manipulate the document objects on the page. Again, this makes sense for a browser where we have a DOM and it does not make sense for node where we do not have a DOM. In node we have something kind of similar to document called process. This provides various methods and properties for the node process that is running (e.g. exit() allows us to exit the current node process).

```
[> process.exit( )
```

Why Should You Use Node.js?

The node.js skillset is in high demand with companies like Netflix, LinkedIn and Paypal all using node in production. Node is also useful for developers anywhere on the stack i.e. front end developers and back end engineers.

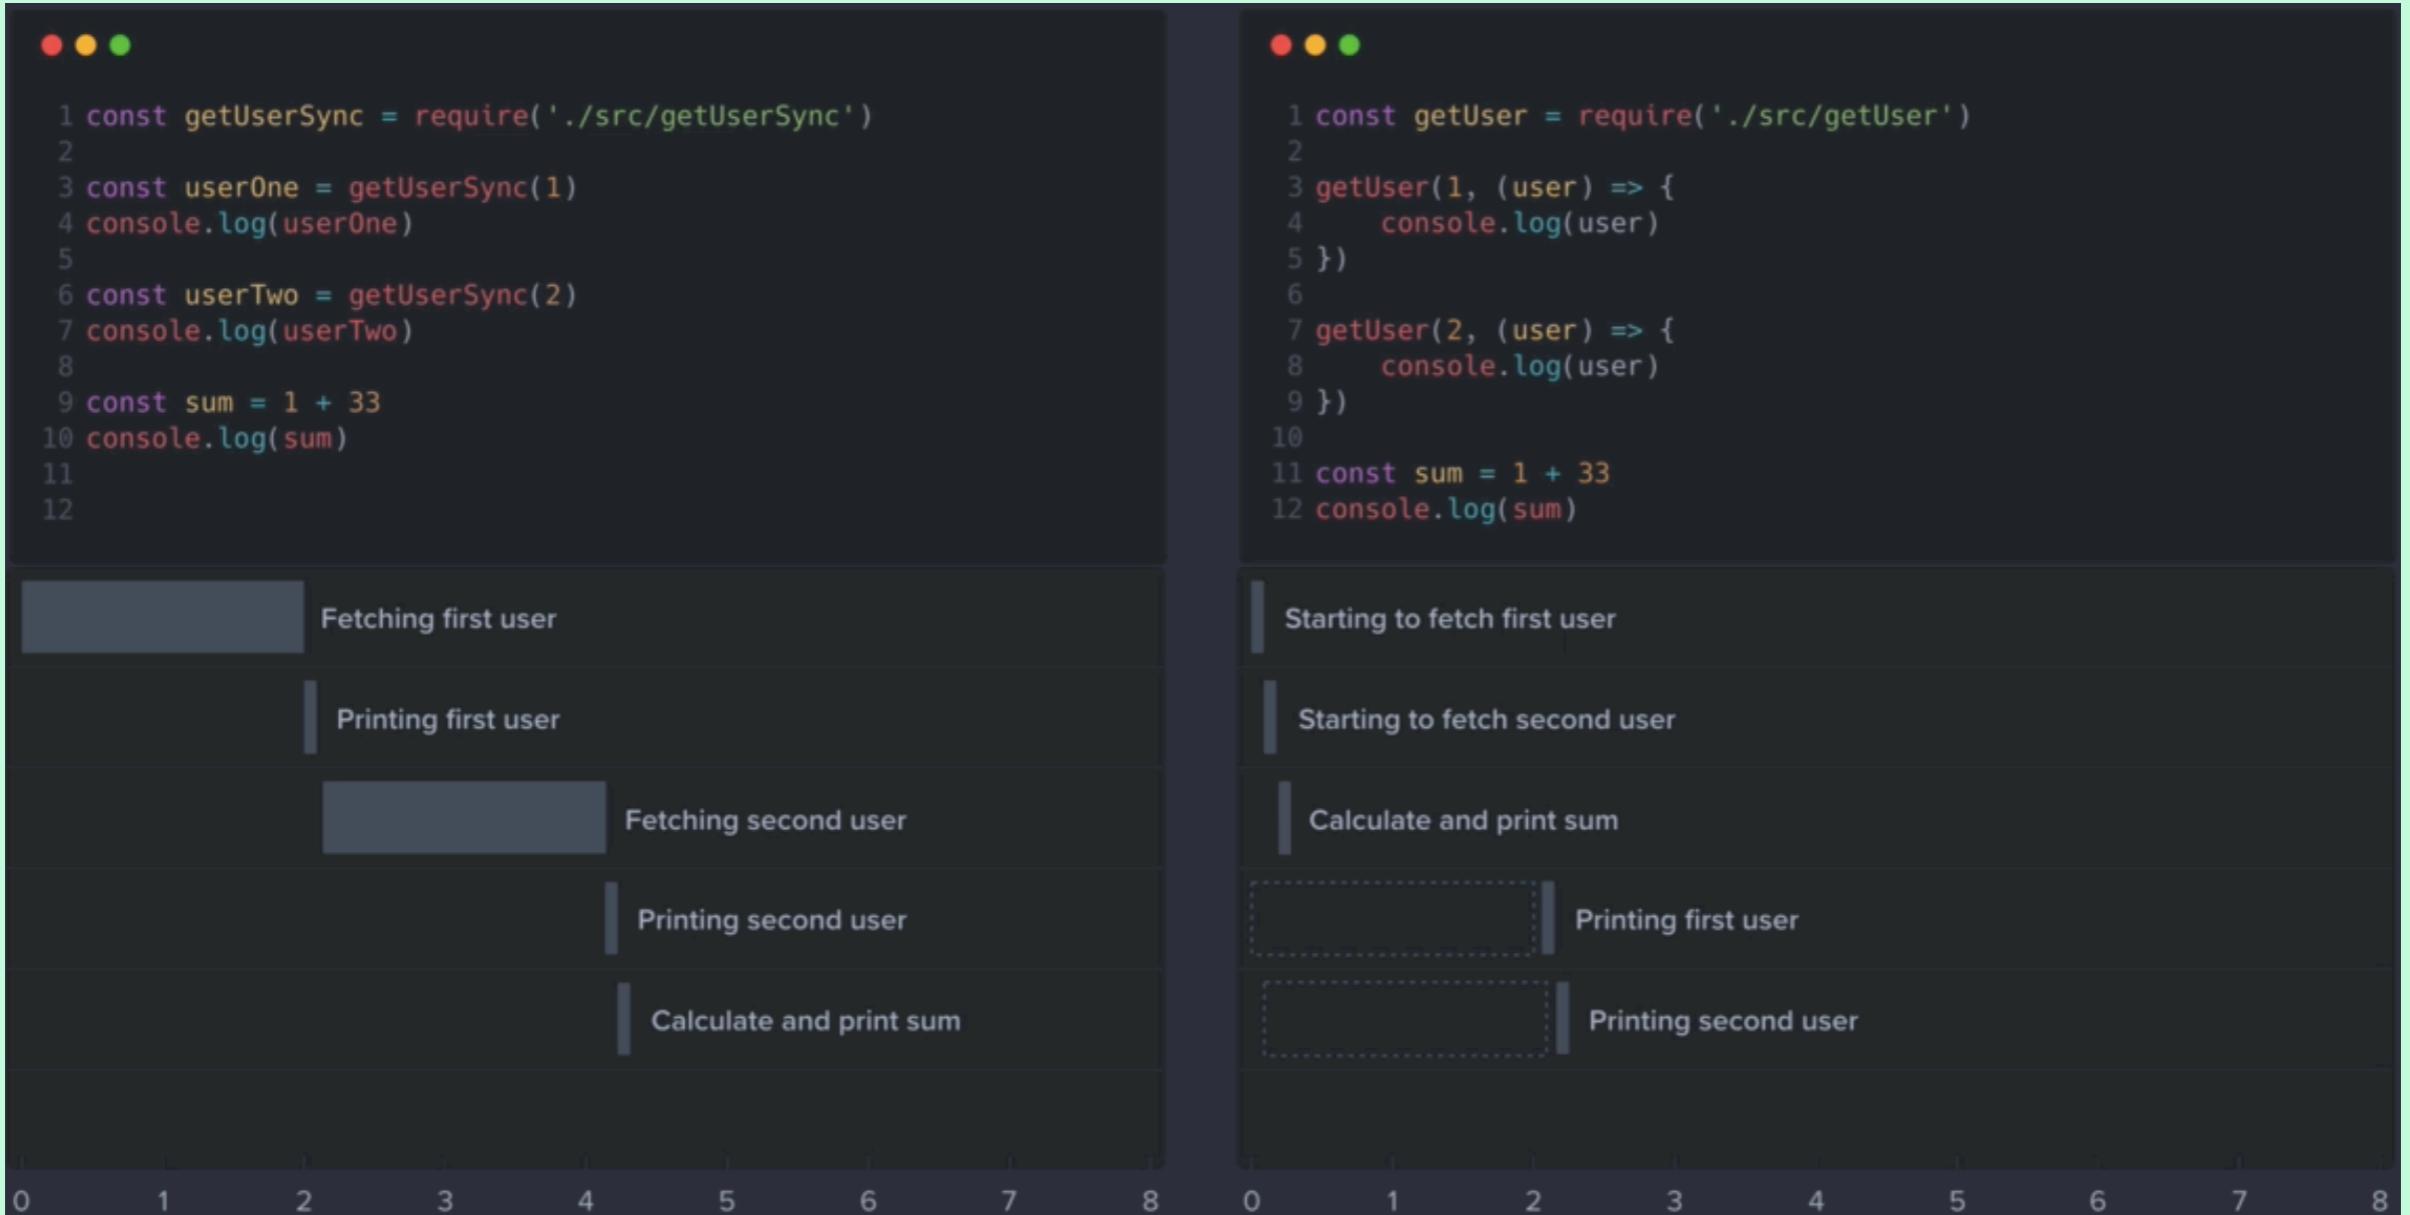
"Node.js uses an event-driven, non-blocking I/O model that makes it lightweight and efficient. Node.js package ecosystem, npm, is the largest ecosystem of open source libraries in the world."

I/O stands for input/output and our node application is going to use I/O anytime it is trying to communicate with the outside world. So if our node app needs to communicate with a machine its running on, that is an I/O for example, reading some data from a file on the file system. If our node app needs to communicate with some other server that is also a I/O for example, querying a database to fetch some records for a given user which is a I/O operation. I/O operations take time. With node.js we get non-blocking I/O which means that while our node application is waiting for a response, it can continue to process other code and make other requests (i.e. Asynchronous).

Below is diagram representation of two I/O (one blocking and one non-blocking) attempting to do the same thing of taking in a user id and go off to a database to fetch the user profile and print the information to the console for two different users and then finally run a calculation of adding two numbers and print that sum to the console. The timeline is going to help us compare the relative operating speeds of the blocking and non-blocking I/O codes.

Blocking I/O

Non-Blocking I/O



The non-blocking I/O is twice as fast at running the code compared to the blocking I/O. This is because we were able to overlap the part of our application that took the longest, which was waiting for the I/O to finish and continue to complete other tasks in the code. This is what makes node.js ideal to develop with.

The node package manager (npm) was also installed on our machine when installing node. This package manager allows us to download open source pre-written packages from the web that we can use inside of our applications as dependencies. We can find all sorts of available packages on <https://www.npmjs.com> and we can use the terminal to install these within our application directory. These libraries and packages makes our life much easier when developing with node.js and we will make heavy use of npm packages which real developers do when they are building out their node applications.

Your First Node.js Script

Using Visual Studio Code, we can create a new file within a directory and name the file with an extension of .js at the end. For example, we can create a file called hello.js which will be our first node.js script file. This is a JavaScript file similar to what we would write if we were to run it in the browser. Our first line of code will use the console.log() to print a message to the console for anyone who runs this script.

hello.js

```
1. console.log('Hello world!');
```

We can use the node documentation to view all the different modules that we can work with based on the version of node we are running on our machine (<https://nodejs.org/en/docs/>).

We can run the following command in the terminal (*VS Code has an integrated terminal we can also use*):

```
:~directoryPathContainingNodeScript$ node hello.js
```

We must ensure within the terminal that we are within the directory path of the file we want to run.

Section 2

NODE.JS MODULE SYSTEM

The Node.js Module System

The node.js module system is the most fundamental feature of node. This is going to allow us to take advantage of all of the interesting things that node provides.

There are all sorts of modules available in node.js and some, for example the Console, are available to us globally. This means that we do not need to do anything special in order to use them. We access console in order to use it. However, other modules require us to actually load them in before they can be used in our scripts for example, the File System module.

The File System module allows us to access the operating system's file system. We will be able to read and write, append files and check if a given file or directory exists and many more interesting features.

We will learn more about this module system.

Importing Node.js Core Modules

We can create a new node app directory and name this anything we want and create a .js file within it. In the below example we are going to explore the File System module function called `writeFileSync()`. This function takes in two arguments both being strings. The first argument is name of the file and the second is the data/content to write in that file.

notes-app > app.js:

1. `fs.writeFileSync('notes.txt', 'This file was created by node.js')`

We can use terminal to execute `node` to run our script. However, we will run into a problem as this is not going to work and throw an error message of `fs` is not defined. The Files System module has to be required by us in our script we are using it in. Before we use the File System module, we need to import it in our script.

This is done using the `require` function that `node` provides which is at the very core of the module system. We use `require` to load in other things whether it's a core node module, another file we created or a npm module we chosen to install and use in our projects. The `require` function takes in a single string.

notes-app > app.js:

1. `const fs = require('fs')`
2. `fs.writeFileSync('notes.txt', 'This file was created by node.js')`

When we are looking to load in a core node module, we just provide the name of the module. For the File System module this is 'fs'. The require function returns all of the things from that module and we have to store this in a variable.

In the above example, we created a const called fs (this could be named anything we want) and its value is going to come from the return value of calling require function (i.e. the node module). We then call on the fs variable followed by the method of writeFileSync to write some text to the notes.txt file.

To run this script we can use the terminal or VS Code integrated terminal and cd into the directory containing our node script. We would call on the node command followed by the name of the script we wish to run:

```
:~/.../notes-app$ node app.js
```

When we run the code we will see that we are brought back to the command prompt asking us to do something else, but we will notice in the tree view for our directory, we have a brand new file called notes.txt which is the file created by the script that is now sitting along side of our app.js file.

The writeFile and writeFileSync methods are responsible for writing data to a file and if the file does not exist it will be created. If the file does exist, its text content will be overridden with the new text content.

We have now learnt how to load in our first core node module and we used it to do something interesting. To know how to load in a core module and what variable naming convention to use or even if the module requires to be loaded in at all, the node documentation is going to be our best friend. We can see the File System module documentation on the following link. Scrolling past the table of content, one of the first things we are going to see is the code to use in order to load the core module i.e. what node calls the module, the common used variable name and whether it needs to be imported in:

<https://nodejs.org/api/fs.html>

Important Note: we do not need to stick to the common variable name; however, it is advisable/best practice to do so, as the naming convention allows developers who are working with our project know which core modules they are working with when importing the modules.

CHALLENGE:

Append a message to the existing notes.txt file using appendFileSync.

1. Use appendFileSync to append to the file
2. Run the script
3. Check your work by opening the file and viewing the appended text.

SOLUTION:

notes-app > app.js:

1. const fs = require('fs')
2. fs.appendFileSync('notes.txt', '- This is a appended text.')

```
:~.../notes-app$ node app.js
```

Importing Your Own Files

When we pass a file to node, only that file executes. This would mean that we would need to put all of our code in a single file if we want to run the code. However, this is not ideal especially if our application grows larger and more complex. This also makes it difficult to expand and/or maintain our application code.

Therefore, we would ideally create our project using multiple files so that we can stay organised and our application is modular. For example, we can define a function in one file and then require it in another file in order to use that function.

To import our own files into node we have to require it in order for the file to get loaded in when we run node. We would continue to use the require function and pass in a single string value. The string value we pass in is a relative path from the file we are loading it in from.

Below is an example of loading in our own file called utils.js which has a single function that logs to the console the name of the file. We will import this file into our app.js and call on this function:

notes-app > utils.js:

```
1. console.log('utils.js')
```

notes-app > app.js:

```
1. require('./utils.js')
2. const name = 'John Doe'
3. console.log(name)
```

The ./ will take us to the relative path from app.js which is the notes-app directory. We can then select the utils.js file that is located within that directory in order to require the file and import it into our app.js file.

When we load in a JavaScript file, it will execute that file when we run our node command on our app.js script. Therefore in the console we should see printed:

utils.js

John Doe

The utils.js file will be printed first because this code runs as soon as we require/load it in our app.js file and then the name variable is printed as it is further down in the code. We now have a simple application that

takes advantage of multiple files.

We can take the above example further and try to define variables within our utils.js file and try to use the variable within our app.js file:

notes-app > utils.js:

1. const name = 'John Doe'

notes-app > app.js:

1. require('./utils.js')
2. console.log(name)

If we were to run app.js script, we will see the error of '*name is not defined*'. This is one very important aspect of the node module system. All of our files, which we can refer to as modules, have their own scope. Therefore, app.js has its own scope with its own variables and utils.js has its own scope with its own variables. The app.js file cannot access the variables from utils.js even though it was loaded in with require.

To share the variables and functions within the utils.js, we need to explicitly export all of the code within the file to be able to share with the outside world i.e. share outside of its own scope. To do this we take advantage of another aspect of the module system which is called module.export. This is where we can define all of the things the file should share with other files.

Using the above example and taking it one step further, is to try and define variables within our utils.js file and export that variable to use it within our app.js file:

notes-app > utils.js:

1. const name = 'John Doe'
2. module.exports = name

notes-app > app.js:

1. const firstName = require('./utils.js')
2. console.log(firstName)

In the above example we have now exported one variable which is a string. We will later learn how to share an object which has a bunch of different methods on it allowing us to export a whole bunch of things. So in utils.js we defined a variable and we exported that variable and other files can now take advantage of the variable.

Whatever we assign to module.exports is available as the return value from when we require the file. Therefore, when we require utils.js in our app.js, that return value is whatever we assigned in the utils.js, which was the string '*John Doe*' that is stored on the name variable.

Within our app.js file we can create a variable and call it whatever we want including '*name*' - variables

within the app.js file is independent from variables from other files because they have different scopes. We assigned this variable value to the require function, which returns the module.export variable value assigned from the utils.js file.

Our application is now back to a working state whereby we can run the app.js script and have the console log the name variable from the utils.js file.

```
:~/notes-app$ node app.js  
John Doe
```

We can also import functions in the same fashion as we did with variables:

notes-app > utils.js:

1. const add = function(a, b) { a + b }
2. module.exports = add

notes-app > app.js:

1. const add = require('./utils.js')
2. const sum = add(1, 2)
2. console.log(sum)

The sum function will return 3 as the value because $1 + 2 = 3$.

```
:~/notes-app$ node app.js  
3
```

CHALLENGE:

Define and use a function in a new file.

1. Create a new file called notes.js
2. Create getNotes function that returns "Your notes..."
3. Export getNotes function
4. From app.js load in and call the function printing the message to the console

SOLUTION:

notes-app > notes.js:

1. const getNotes = function() { return 'Your notes...' }
2. module.exports = getNotes

notes-app > app.js:

1. const getNotes = require('./notes.js')
2. const msg = getNotes()
3. console.log(msg)

Importing NPM Modules

We can use the Node Module System to load in npm packages which will allow us to take advantage of all of the npm packages inside of our node applications. NPM modules allows us to install code written by other developers so that we do not need to reinvent/recreate the wheel from scratch.

There are always things that every applications out there needs to do such as validating data such as emails and maybe even sending an email. These are core functionality which is not specific to what our application does for our users. So if we use npm modules to solve common problems (which is the standard in the node community) then we can spend our development time focusing on features that makes our app unique.

When we installed node on our machine, we also got the npm program installed on our machine. This gives us access to everything at <https://www.npmjs.com>. We can use npm terminal commands such as the one below which shows the npm version installed on our machine:

```
:~/notes-app$ npm -v
```

Before we can use npm modules within our scripts, we have to take two very important steps.

1. We have to initialise npm in our project
2. We have to install all of the modules we actually want to use

To initialise npm in our project, we need to run a single command in the project root directory:

```
:~/notes-app$ npm init
```

This command is going to initialise npm in our project and create a single configuration file that we can use to manage all of the dependencies from the npm website that we want to install. The above command will ask a few questions to configure the configuration file such as package name (i.e. the folder/project name), version, description, etc. We can type in our own values to overwrite the default values to the questions. Once we complete all of the questions and execute the command this will create a package.json file in the root directory of our projects. This is the configuration file which contains all the npm module dependencies and scripts for our project.

We can go onto <https://www.npmjs.com> to look for packages/modules we wish to use in our project. Once we find a package that we want to install we can use the npm terminal command within our project directory. The below example demonstrates installing the validator module package:

```
:~/notes-app$ npm install validator@11.1.0  
:~/notes-app$ npm i validator@11.1.0
```

When we run the command, we will notice that two things have occurred. Firstly, we would now have a package-lock.json file and secondly, we have a new node_modules directory (folder) in our project root.

The `node_modules` is a folder that contains all of the code for the dependencies that we have installed. In the above example, this would have a subfolder called `validator` which contains all of the code for that dependency package we installed using npm. We should never go into the `node_modules` folder to manually change these files as this is simply a package management. When working with `node_modules`, it is going to get generated and edited when we run the `npm install` command. The npm maintains this directory.

The `package-lock.json` file contains extra information for npm to make npm a bit faster and secure and again is another file that we do not modify. This file lists out the exact versions of all of our dependencies as well as where they were fetched from and the hash masking sure that we are getting the exact code that we got previously, if we were to install a dependency again. This file will be maintained by npm.

When we ran the above command, the package is also added as a dependency within our `package.json` file which lists the package name and the version installed.

Now that we have our dependency installed, we can now load in the dependencies within our application's files using `require` and take advantages of the functionality it provides.

notes-app > app.js:

```
1. const validator = require('validator')
```

To load in npm modules we list out the package names similar to node core modules. Require is going to return all of the stuff that the validator package provides us. We can create a variable and assign it to the require function which gets its contents from the package it is returning.

When it comes to figuring out how to use a given npm package, this is when we have to turn to the documentation of that package to understand how it was intended to be used. The below example provides an example of using the validator package and one of its many functions in our app.js file:

notes-app > app.js:

1. const validator = require('validator')
2. console.log(validator.isEmail('test@email.com'))

```
:~/notes-app$ node app.js  
true
```

Is we run our application, the validator isEmail function will validate whether the given string is actually an email. In the above example, this equated to true.

As mentioned above, the node_modules is not a file that we should be manually editing. This is because when we use the npm command again, our edits are going to be overwritten. We can actually recreate this

directory from scratch using npm based off of the contents of package.json and package-lock.json.

The node_modules folder is necessary for our app to run if we are using npm modules as dependencies for our application code. To install the node_module folder with all the dependencies, we can run a simple npm command:

```
:~/.../notes-app$ npm install
```

This is going to look at the package.json and package-lock.json to determine which dependencies and versions our application is using in order to recreate the node_module folder from scratch based on the contents of these two .json files (*the .json extension stands for JSON Object Notation*).

This is useful for when sharing your code as the node_modules folder can get really huge in file size and it would be easier to regenerate the folder using the npm install command using the two package files. Our application will be back to its working state as we would have all the necessary files in order to make our application work again.

CHALLENGE:

Install and use the chalk library.

1. Install version 2.4.1 of chalk
2. Load chalk into app.js
3. Use it to print the string "Success!" to the console in green
4. Test your work
5. Bonus: use chalk documentation to play around with other styles e.g. make text bold and inverse.

SOLUTION:

```
:~/notes-app$ npm install chalk@2.4.1
```

```
notes-app > app.js:
```

1. const chalk = require('chalk')
2. console.log(chalk.green('Success!'))
3. console.log(chalk.blue.bgYellow.bold.inverse('Hello World!'))

```
:~/notes-app$ node app.js
```

Global NPM Modules and nodemon

Installing Global npm module packages will allow us to get new commands that we can execute from the terminal. So far we have learnt how to install npm packages locally, this is where we install the dependencies explicitly into our project directory and these appear in the package.json and package-lock.json file as dependencies.

We can install a npm package called nodemon globally which is a nice utility when working with node. This is going to allow us to run our application and automatically restart the app whenever the app code changes i.e. whenever we save any changes to our app code file. This would mean that we would not have to constantly switch to the terminal and rerun the same command over and over again to test our code.

The documentation for nodemon can be found on <https://www.npmjs.com/package/nodemon> which we can read to understand more about the package and how to use it.

When installing packages globally, the command is exactly the same as if we were installing it locally, however, there is one slight difference - we add the -g flag to indicate installing the package to be global.

```
:~$ npm install nodemon@1.19.1 -g  
:~$ sudo npm install nodemon@1.19.1 -g
```

On linux and MacOS, we can use the sudo in the command to install the package as an administrator to avoid errors in the command.

It is advisable to install packages locally rather than globally as you can tend to run into errors installing packages globally depending on the OS environment. Furthermore, packages are updated regularly and this would mean older and newer projects would have different dependency packages at different versions that are used in development which could also create problems for your development environment and your code running especially if there are syntax differences or deprecation in the dependency packages used. Therefore, it is more advisable to either install packages locally or have multiple virtualised development environments whereby we can install packages globally.

To check that we installed nodemon correctly gobally, we can run the following command:

```
:$ nodemon -v
```

If installed locally we can use the following command to check the version installed:

```
:~/notes-app$ ./node_modules/nodemon/bin/nodemon.js -v
```

To start using the nodemon to automatically run our application script we would run the following command:

```
:~/notes-app$ nodemon app.js
```

We would use the following command if we were running it locally:

```
:~/notes-app$ ./node_modules/nodemon/bin/nodemon.js app.js
```

We will notice that when we run nodemon and our script has run, it is not bringing us back to the command prompt even though our application has finished, this is because the nodemon process is still executing. If we make any changes it will automatically run our code when we save the file i.e. the process is now listening for any saved changes to our file.

The nodemon package is a great utility to improve our overall developer experience when developing node applications.

To terminate the nodemon process, within the terminal running nodemon, we can press control and c on our keyboard and this will end the process from running.

Section 3

FILE SYSTEM & COMMAND LINE INTERFACE ARGUMENTS

Getting Input From Users

Input from the user is essential to create anything meaningful. This allows interaction between the user and the application.

We can get input from the user either from a command line argument or from a client such as a browser. We are going to focus on the fundamentals of getting input from the user with command line arguments.

We can run our app using the node command in the terminal. However, we can also pass in additional information that our application can choose to use to do something dynamic for example print a greeting with the name inside of the message:

```
:~/.../notes-app$ node app.js "John Doe"
```

The above command will run our script as normal, but does not do anything with the additional information

we passed in. The question follows, where exactly do we access those command line arguments? The answer is that on the global process object there is a property that allows us to access all of the command line arguments passed into our app.

Process is an object that contains many methods and properties and the property we are interested in is the `argv` (*argument vector*) property. This property is an array that contains all of the arguments provided.

notes-app > app.js:

1. `console.log(process.argv)`

If we now rerun our script with John Doe appended, we should now receive a new array of the command line arguments:

```
:~/notes-app$ node app.js "John Doe"  
[..., ..., 'John Doe']
```

There are always 2 strings provided by the `argv` property, the first one is the path to the node.js executable on our machine and the second is the path to our `app.js` file. These paths are going to be slightly different depending on where the file lives on the machine. The third value and all values after are the arguments the user provides which we can take advantage of the string in our application to do something meaningful.

We can extract the individual values from the array by using JavaScript bracket notation to retrieve the value from its index. JavaScript uses zero indexing whereby 0 is the first item within the array.

notes-app > app.js:

1. console.log(process.argv[2])

```
:~/notes-app$ node app.js "John Doe"
```

```
John Doe
```

We can use the command line argument to check for the value and then perform a certain code in our application. For example, we can take the first argument as a command and look at the value passed in such as add and then run a particular code to add a note:

notes-app > app.js:

```
1. const command = process.argv[2]
2. if (command === 'add') { console.log('Adding note!')
3. } else if (command === 'remove') { console.log('Adding note!') }
```

```
:~/notes-app$ node app.js add
```

```
Adding note!
```

We now have a way to get the users input and perform some form of meaningful action. Now that we can get the command from the user, we need to get a command line options so that the user can provide additional information for example the note title or note body they wish to add/remove.

```
:~/notes-app$ node app.js add --title="This is my title"
```

If we were to view the `process.argv` we will notice that this command line option will not be parsed. We would see in the console `--title="This is my title"` printed. We would have to figure out how to parse the title and then figure out whether the `--title` option was provided and if so, we have to get that value.

If we wrote our own code for parsing the string, we would have to write tests and maintain this code and this code does not do anything unique to our application. This would be a best case scenario where we would look for a npm package that can take care of parsing for us.

We have now found a way of our users interacting with our application and getting their input via the command line arguments.

Argument Parsing with Yargs

Node does not provide any argument parsing and it is a very bare bones utility allowing us to access the

raw elements. So when we pass in an argument such as `--title="This is my title"` option, it was not parsed in a way that is particularly useful. We would need to write some code to extract the key value pair.

Most node.js applications use command line arguments in some way and there are many of great npm packages that make it easy to setup our commands and options as we want. Yargs is a widely used and tested utility package that will help us to parse command line arguments (<https://www.npmjs.com/package/yargs>). To install Yargs in our project we can run the following command:

```
:~/notes-app$ npm install yargs@13.3.0
```

Once installed, we can use yargs in our project to parse command line arguments. To use yargs we first need to require the npm package. When using require, it is good practice to use a pattern i.e. require core modules first, then npm packages and finally our own files.

notes-app > app.js:

1. `const yargs = require('yargs')`
2. `console.log(yargs.argv)`

We can compare and contrast the yargs command with node's `process.argv` command by running the `app.js` script and passing in arguments, in order to see the difference between the two approaches when parsing command line arguments.

```
:~/notes-app$ node app.js  
{ _: [], '$0': 'app.js' }
```

We would notice by running the first example, yargs almost provides nothing compared to the default. In the yargs output we have an object with two properties: the first is an underscore property which will get populated with various arguments and the second is the \$0 property which has the value of the file name we executed.

```
:~/notes-app$ node app.js add --title="This is my title"  
{ _: ['add'], title: 'This is my title', '$0': 'app.js' }
```

Running the second example, we can see yargs provides us with a much more useful parsed object. Our commands appears under the underscore property and we then have an actual options property on this object that contains the string value (i.e. `title: 'This is my title'`). Yargs has done the heavy lifting of parsing our options and putting them on the object so that they are easy to access. We can access the title property to do something with it such as creating a new note and save it to our data store.

This is the very bare bones way of using yargs, since we have not configured it to do anything special. By default we get some very useful functionality behaviour. If we run the below command, we can see the version of yargs (by default this is 1.0.0)

```
:~/notes-app$ node app.js --v  
1.0.0
```

We can change the version by accessing `yargs` and calling the `version` method. This takes in a string as its one and only argument as the new version number. This allows us to specify a different version number than the default 1.0.0 version.

notes-app > app.js:

1. `const yargs = require('yargs')`
2. `yargs.version('1.1.0')`

```
:~/notes-app$ node app.js --v  
1.1.0
```

`Yargs` comes with a lot of useful parsing and features of its own but it can be configured to do anything we want, for example we can setup distinct commands, display help options for the commands, create options to pass in data, etc. all of which is supported by `yargs`.

The `yargs command()` function takes in an object, which is the `options` object, where we can customise how this command should work. The first command parameter is the name of the command. The second `describe` parameter allows us to add a description of what the command is suppose to do. We can setup a

third handler property which is the code that is actually going to run when someone uses the command described in the first parameter. Below is a example yargs command code:

notes-app > app.js:

```
1. const yargs = require('yargs')
2. yargs.command{
3.   command: 'add',
4.   describe: 'Add a new note.',
5.   handler: function(){
6.     console.log('Adding a new note!')
7.   }
8. })
9. console.log(yargs.argv)
```

If we use the yargs --help command, this will now add a new Commands line displaying all of the different commands we have setup in yargs along with the help description of the commands.

```
:~/notes-app$ node app.js --help
```

Commands:

```
  app.js add  Add a new note.
```

Options: ...

We can now also use the add command within our command line argument:

```
:~/notes-app$ node app.js add
```

```
Adding a new note!
```

Now that we know how to use the yargs utility package to create our command line argument options we now need a way to configure options for those commands. To setup a configuration option, there is another property we can add to our configuration object that we pass to our yargs command function called builder. The builder value is an object where we can define all of the options we want the given command to support. Each option value is also a object where we can customise how the option works:

notes-app > app.js:

```
1. const yargs = require('yargs')
2. yargs.command({
3.   command: 'add',
4.   describe: 'Add a new note.',
5.   builder: {
6.     title: { describe: 'Note title.', demandOption:true, type: 'string' }
7.   },
8.   handler: function(argv) {
9.     console.log('Adding a new note!', argv)
10.  }
11. })
```

12. `yargs.parse()`

The option property value takes in a describe property which describes the option. With our option defined we can now access that in the command handler and we get access via the first argument provided in our function. If we now run the following command in terminal:

```
:~/notes-app$ node app.js add --title="This is my title"  
Adding a new note! { _: ['add'], title: 'This is my title', '$0': 'app.js' }
```

We will notice that we get the console logging the text along with yarg's version of argv inside of the handler. We also have access to the options and its values for example argv.title which we can now use to actually define a note.

One thing to note is that the options are not required and we can run our terminal command without using any of the options and things will continue to work. It is up to us as the developer to decide whether or not a given option should be required. If we want a option to be required we can set this up as well using the demandOption property. By default this property is set to false. If we set this to true, this would mean that we must provide the option in order for the command to function correctly.

If we now try to run the command without the required option this will print an error message of missing required argument: ... listing out the missing command options. It will also display all the available options

and which ones are required. If we run the command with the required option but without a value:

```
:~/notes-app$ node app.js add --title  
Adding a new note! { _: ['add'], title: true, '$0': 'app.js' }
```

We will notice that the value for title will default to a boolean value. To set the argv --title option property to be a string value, we can enforce that by setting the type property on our option object configuration. We can set this to one of the supported types as a value for example boolean, string, number, array, etc.

If we now run the following command without a value, at least the value for title will now be a string, even if this is an empty string value.

```
:~/notes-app$ node app.js add --title  
Adding a new note! { _: ['add'], title: "", '$0': 'app.js' }
```

In order to view the output of yarg we need to `console.log(yargs.argv)` at the end of our code else we would not see any output. If we do not need to access the `yargs.argv`, which is less than ideal, we can parse it using the `yargs.parse()` function.

We can have access to our options property on the argv for example within the handler we can write:

```
9.           console.log('Title: ' + argv.title)
```

This will print whatever value the user passes in for the title option that was setup.

We now know how to setup yargs to create our own commands and options for the user to interact with our node application via the command line arguments and we can use this utility to parse the arguments in order for our application to do something meaningful.

CHALLENGE:

Add new command using Yargs.

1. Setup command to support “read” command (print placeholder message for now)
2. Test your work by running command and ensure correct output in the terminal

SOLUTION:

notes-app > app.js:

```
1. const yargs = require('yargs')
2. yargs.command({
3.   command: 'read',
4.   describe: 'Read a note.',
5.   handler: function() {
6.     console.log('Reading a note!')
7.   }
8. })
```

9. console.log(yargs.argv)

```
:~/notes-app$ node app.js read  
Reading a note!
```

CHALLENGE:

Add an option to Yargs.

1. Setup a body option for the add command
2. Configure the option with a description, make it required and set the type to string
3. Log the body value in the handler function
4. Test your work!

SOLUTION:

notes-app > app.js:

1. const yargs = require('yargs')
2. yargs.command({
3. command: 'add',
4. describe: 'Add a note.',
5. builder: {

```
        body: {  
            describe: 'Note body',  
            demandOption: true,  
            type: 'string'  
        }  
    },  
5.    handler: function(argv) {  
6.        console.log('Body: ' + argv.body)  
7.    }  
8. })  
9. yargs.parse( )
```

```
:~/notes-app$ node app.js add --body='This is my body.'  
Body: This is my body.
```

Storing Data with JSON

We can use the `fs` core module to save data to the file system. We can save the data in the JSON data format which structures the data as JavaScript Objects. JSON has native support in JavaScript and is very easy to learn since it is essentially arrays and objects with various properties. We can model our data in JSON and store the data in a `.json` file type which we can access through our application. If we access to these arrays and objects, then we have an array of items like we would in any database.

The fs core module only knows how to work with string data. In order to take a JavaScript object and store it as a string representation, we would use the JSON.stringify() method:

playground > json.js

```
1. const book = {  
2.   title: 'Ego is the enemy',  
3.   author: 'Ryan Holiday'  
4. }  
5. const bookJSON = JSON.stringify(book)  
6. console.log(bookJSON)
```

The stringify() method takes in an object, array or any value and it returns the JSON string representation. In the above example, we can pass in the book variable object and convert it to a JSON string:

```
:~/playground$ node json.js  
{ "title": "Ego is the enemy", "author": "Ryan Holiday" }
```

We will notice that the JSON string data looks exactly like a JavaScript object with the only difference of all properties name and values wrapped in double quotes. The JSON data is a string and not an object and therefore we cannot access its properties like we would with a JavaScript object i.e. bookJSON.title does not exist unlike the book.title object.

To convert a JSON data back into a JavaScript object we would use the JSON.parse() method:

playground > json.js

...

7. parsedData = JSON.parse(bookJSON)
8. console.log(parsedData.author)

This method takes in the JSON string and converts it into a JavaScript object which we can then have access to the object properties:

```
:~/playground$ node json.js  
Ryan Holiday
```

These are the two core method with working with JSON data. We can integrate this with the file system core module to stare and retrieve data from our application using JSON.

playground > json.js

1. const fs = require('fs')
2. const book = { title: 'Ego is the enemy', author: 'Ryan Holiday' }
3. const bookJSON = JSON.stringify(book)
4. fs.writeFileSync('book-json.json', bookJSON)

We can use the fs.writeFileSync function to create a new file in our file system called book-json which will have the file extension of .json as this will be a JSON file containing JSON data. The second argument we pass is the data to write to the file which is going to be the bookJSON data. If we now run the above code

this should create a new file in the same directory as the script we ran called book-json.json and this file would contain our JSON string data:

```
:~/playground$ node json.js
```

We now have a way to store data with JSON and the file store core module using our node application. This is now a data we can load in at a later time and use within our node application.

To load in a data with the file system core module, we would use the fs.readFileSync method and pass in a single argument, which is a string of the file name we are trying to read:

playground > json.js

1. const fs = require('fs')
2. const dataBuffer = fs.readFileSync('book-json.json')
3. const dataJSON = dataBuffer.toString()
4. const data = JSON.parse(dataJSON)

We can grab the content of the returned value from the fs.readFileSync which is going to contain the contents of the file. We can create a variable that will store the content that came back from the method call. What comes back from the method is not a string, but actually a buffer which is a way for node.js to represent binary data. We can use the toString() method to convert the buffer into a string as we would expect to see the returned contents. We can finally parse the string into a JavaScript Object which we can have access to its properties and values.

CHALLENGE:

Work with JSON and the File System.

1. Load and parse the JSON data

data.json:

1. {"name": "Andrew", "planet": "Earth", "age": 28}

2. Change the name and age property using your info

3. Stringify the changed object and overwrite the original data

4. Test your work by viewing data in the JSON file

SOLUTION:

playground > jsonChallenge.js:

```
1. const fs = require('fs')
2. const dataBuffer = fs.readFileSync('data.json')
3. const dataJSON = dataBuffer.toString()
4. const user = JSON.parse(dataJSON)
5. user.name = "Beth"
6. user.age = 21
7. const userJSON = JSON.stringify(user)
8. fs.writeFileSync('data.json', userJSON)
```

:~.../playground\$ node jsonChallenge.js

ES6 Arrow Functions

To understand the ES6 arrow functions, we first need to compare it with a simple ES5 function. Below is a regular function that takes a number and returns the squares of the given number using the ES5 syntax:

playground > arrow-function.js:

1. `const square = function(x) { return x * x }`
2. `console.log(square(3))`

```
:~/playground$ node arrow-function.js
```

```
9
```

The square of 3 will return 9 ($3 \times 3 = 9$) and therefore is printed to the terminal when we run the script.

Below is the equivalent function but using the new ES6 arrow function syntax:

playground > arrow-function.js:

1. `const square = (x) => { return x * x }`
2. `console.log(square(3))`

```
:~/playground$ node arrow-function.js
```

```
9
```

This too will return 9 in the terminal, and the function continues to operate using the ES6 syntax. The first difference we can see is that we do not use the `function` keyword, we actually start with the argument list.

The next step is to put the arrow function by using the `=>` sign/symbols to create the arrow. Finally, we can setup the curly brackets and define the code block for what the function is suppose to do. Now that we know how to setup a basic arrow function, we can now explore its hidden features that make it a tool worth using:

The first feature is its shorthand syntax. With a lot of the functions we end up writing in JavaScript and node.js they end up being pretty simple. We take in some arguments and then immediately return some result. If our function was just going to have a single statement which returned something, we can put that something right after the arrow function:

playground > arrow-function.js:

1. `const square = (x) => x * x`
2. `console.log(square(3))`

`:~/playground$ node arrow-function.js`

`9`

There is no need for both the curly brackets or the `return` keywords. Whatever we put after the arrow function is going to be implicitly returned. Now, in reality we would not be able to use this shorthand syntax for every arrow function that we create, but we can use it for simpler function which immediately returns some sort of value. If we have longer code blocks such as using `if` statements, we would use the long form version where we setup curly brackets and add as many lines of code as needed.

We can now take this up a gear and look at arrow functions work in the context of methods, so arrow functions as properties on an object. Below is a example of an event object with two properties. The second property calls on a function using the ES5 function syntax:

playground > arrow-function.js:

```
1. const event = {  
2.   name: 'Birthday Party',  
3.   printGuestList: function( ) { console('Guest list for ' + this.name) }  
4. }  
5. event.printGuestList( )
```

```
:~/playground$ node arrow-function.js
```

```
Guest list for Birthday Party
```

The function `printGuestList` will start with a static text of "*Guest list for*" followed by the actual event name i.e. "*Birthday Party*" - we can read the `name` property value. We know that with our methods, our functions as object properties, we have access to the original object via the `this` binding. Therefore, this is a reference to our object, which means we can access properties on the `this` keyword such as `this.name`. We can use the `printGuestList` property method on the event object to print the text to the console.

If we now view the above code but written using the ES6 arrow function syntax, we can view the difference between the two syntax and the second hidden feature provided by ES6 arrow functions:

playground > arrow-function.js:

```
1. const event = {  
2.   name: 'Birthday Party',  
3.   printGuestList: () => { console('Guest list for ' + this.name) }  
4. }  
5. event.printGuestList()
```

```
:~/playground$ node arrow-function.js
```

```
Guest list for Birthday Party
```

The above is the valid syntax but we will notice that if we run the code, we do not get the results that we might be expecting. The arrow function is unable to find the name property and this is because arrow functions do not bind their own “this” value, which means we do not have access to the this keyword as a reference to the object because of the fact we are using the arrow function.

Arrow functions are not well suited for method properties, that are functions, when we want to access “this”. In this case it would be best to use a standard ES5 function. This does not mean we are stuck with the ES5 syntax. There is actually a ES6 method shorthand that allows us to use a shorter more concise syntax while still having access to standard function features like the this binding:

playground > arrow-function.js:

```
1. const event = {  
2.   name: 'Birthday Party',  
3.   printGuestList( ) { console('Guest list for ' + this.name) }  
4. }  
5. event.printGuestList( )
```

```
:~/playground$ node arrow-function.js
```

```
Guest list for undefined
```

We remove the colon and function keyword keyword so that it goes from the method property name right into the arguments list followed by the function body contained in the curly brackets. This is now a standard function which does have the this binding and it is not an arrow function. This is using an alternative syntax available to us when we are setting up methods on objects. If we run the script, we would now have a working code block using this concise syntax.

Why do arrow functions avoid their own this binding? Below is an example of the same function but we now have a guestList property on our event object:

playground > arrow-function.js:

```
1. const event = {  
2.   name: 'Birthday Party',  
3.   guestList: ['Alan', 'Beth', 'Carl'],
```

```
4.     printGuestList( ) {
5.         console('Guest list for ' + this.name),
6.         this.guestList.forEach(function( guest) { guest + ' is attending ' + this.name } )
7.     }
8. }
9. event.printGuestList( )
```

```
:~/playground$ node arrow-function.js
```

```
Guest list for Birthday Party
Alan is attending undefined
Beth is attending undefined
Carl is attending undefined
```

If we now try to print that guestList for each individual guest, we can use the `forEach` method on the array using the `this` binding. The `forEach` method takes in a function which gets called one time for every array item within the list i.e. one for each guest. We get access to the guest via the first argument, which we could name as anything we want. We can use `console.log` function to print a message for each guest and we can access the name using the `guest` argument variable. We can see that things are not working as expected as the event name is `undefined`. This has to do with the `this` binding. Standard functions are going to have their own binding (i.e. the `forEach` function has its own binding), which is a problem because we do not want the `forEach` function to have its own `this` binding. We want to be able to access the `this` binding from the parent function which is `printGuestList()`. In the past there were all sorts of work around for this for

example we would have created a const variable of "that" and setting it equal to "this" within the parent function, so essentially creating a reference that we can access later, and in the sub function we would use the "that" variable, e.g. that.name, to access the object property from the parent which would work:

playground > arrow-function.js:

```
1. const event = {  
2.   name: 'Birthday Party',  
3.   guestList: ['Alan', 'Beth', 'Carl'],  
4.   printGuestList( ){  
5.     that = this,  
6.     console('Guest list for ' + this.name),  
7.     this.guestList.forEach(function( guest) { guest + ' is attending ' + this.name } )  
8.   }  
9. }  
10. event.printGuestList( )
```

```
:~/playground$ node arrow-function.js
```

```
Guest list for Birthday Party
```

```
Alan is attending Birthday Party
```

```
Beth is attending Birthday Party
```

```
Carl is attending Birthday Party
```

Although the code will work using the workaround; however, this is not ideal and this is where ES6 arrow functions solves this problem as seen below:

playground > arrow-function.js:

```
1. const event = {  
2.   name: 'Birthday Party',  
3.   guestList: ['Alan', 'Beth', 'Carl'],  
4.   printGuestList( ) {  
5.     console('Guest list for ' + this.name),  
6.     this.guestList.forEach(( guest) => { guest + ' is attending ' + this.name } )  
7.   }  
8. }  
9. event.printGuestList( )
```

```
:~/playground$ node arrow-function.js
```

```
Guest list for Birthday Party  
Alan is attending Birthday Party  
Beth is attending Birthday Party  
Carl is attending Birthday Party
```

The solution is turning the `ForEach` regular function into an arrow function. If we now run the script, we will see the code block running as expected. Arrow functions do not bind their own "`this`" value. They access the "`this`" value in the context which they are created, which in the above case is inside of the `printGuestList`

function. This means that we have access to the `this.name` which is pointing to the property up above within the object.

To conclude, there are three key aspects we have learnt:

1. There is an alternative syntax using ES6 arrow function
2. With arrow functions we have a shorthand syntax for returning immediate values
3. Arrow functions do not bind their own “this” value

Therefore, arrow functions are poor candidates for object methods (i.e. functions on an object) but they are great candidates for everything else. Moving forward, we are never going to use a function where we have the `function` keyword. We are either going to use arrow functions or when necessary we will use the ES6 method definition syntax. This understanding will make us better JavaScript developers.

Section 4

DEBUGGING NODE.JS

Debugging Node.js

When we run our program, things do not always work as expected. We may introduce a typo into our code causing the program to crash or we may misuse a function causing strange and unexpected behaviour.

The goal in this section is to have the tools needed to actually fix these problems. Therefore, debugging and inspecting the application as it is running is an important skill so that we can figure out exactly where the program is going wrong and why. This is going to give the information needed to fix the mistake in the code and get back to a working application.

As we make mistakes it is only going to get easier to fix them over time. What may have taken 10 minutes to debug will only take 10 or 20 seconds to fix. You will notice and start to see a lot of the same issues and patterns over time. The real goal is to recover from the mistake quickly to get back on track and be productive.

We will go through debugging strategies as well as exploring debugging tools provided by node to make it much easier to fix our broken application. We will explore the node debugger which integrates directly with the chrome developer tools, providing us with a nice GUI for debugging our backend applications.

In general, there are two types of errors we can run into when working with node.js and these are:

1. Explicit error message e.g. typos in the class
2. Logical errors i.e. no syntax error or error message

With logical errors there will be no error message, and so we will need to first figure out where it is and then secondly adjust to make the application code actually work the way we envision it.

We can use the same debugging strategies for both type of errors. Firstly, we can explore the basic tools that node provides:

1. The `console.log` command

This allows us to view the value of variables when we run a script which allows us to figure out what is happening with our application. This is the most basic tool available to us for debugging our application.

2. The node debugger

This is the node's built in debugging tool which integrates with V8 and the Chrome browser. The debugger just like `console.log` needs to be added at a specific point in our application code. Using the keyword `debugger` is going to stop the application at a point in time whereby we can use the developer tools to

look at any values we want. This gives us more flexibility, because unlike `console.log` where we have to log every variable we want to see, with `debugger` we can just add the `debugger` statement one time and once we get to the dev tools, we can view everything of our application from where it stops at the debugger.

When we have debuggers in our application, they are not going to pause the program by default, we have to run `node` with a special option to get that done. This special option is called `inspect`:

```
:~/notes-app$ node inspect app.js add --title="t" --body="b"  
< Debugger listening on ws://127.0.0.1:9229/XXXXXXXXXXXXXX
```

We would now get a much different output. Node is letting us know that a debugger is now up and running and which url it is listening on. We can go into chrome following the url to actually inspect our application. We can go into chrome and navigate to <chrome://inspect> which will bring up the Chrome debugger which is using the built in V8 debugger tools. Chrome is the only browser we can inspect our node applications because we require the V8 engine's debugger tools.

Under Remote target we should see two targets i.e. one at the port and another at the shorthand localhost (both are the exact same process). If you see nothing under Remote Target, it is likely you have a misconfiguration. If we click on the configure button, we should see two values:

Localhost:9229 and 127.0.0.1:9229

If we do not have these values, we can add them her into the configuration and should be able to see our

two targets. From here we can actually inspect our application and we can pause at that point in time that we put the debugger statement and view all of our application variables and values. In the Target, we would click on the inspect link which will open a new instance/window of the debugger tools for debugging our Node.js application.

The contents and sources tab are the two most important tabs that we would be using to debug our application. Within the sources tab we have in the middle our file code, exactly how we have it inside of our text editor with one very important difference; we have a function that wraps our whole node.js script.

In other programming languages we might have a main function that we define. When we run the program, that is the function that starts up our application. With Node.js we do not have that as it runs our entire script from top to bottom. Therefore, to make node run our program through V8, our code is wrapped in this function which can be considered as our main function. This function provides to our application a few different values we can access such as exports, require, module, __filename and __dirname variables.

On the left hand side we can add our project folder to the developer tools by clicking on the +Add folder to workspace button under the Filesystem. This will change the middle screen slightly as we no longer will have the wrapper function.

If we press the escape key we can toggle the console at the bottom of the debugger to have both source

and console on the same window. We can use the console to run various commands. It is important to note that at this point in time, not a single line of code in our script has executed. Line one of our code is highlighted in blue which indicates the line that the debugger is paused at. When it is paused at a line, it has not executed that line.

On the right hand side we have a lot of information about where our program currently is at. Information such as: the current call stack, the scope we currently have access to and above tools to work through our application code. We have to manually tell the node debugger tool to continue running our application. The play button will run our application until it is told otherwise. One of the things that is going to pause the application is the debugger statement which we have in our application code. This will open the file that has the debugger statement, if we placed it in another file. At this point in time we can really start to debug our application using the right hand side debugger tools.

The call stack will show us where in the program we are at i.e. the line number of the specific node.js file along with the function called. Within the scope section, we can actually see access to the variables we have in scope. We can dive deeper into those values by running statements from the console e.g. grabbing one of the notes value from the array. We can dump the values in the console and inspect the individual properties to check that the value is as expected it to be.

The debugger statement allows us to dive deeper into the current state of our application and we can

figure out what is not working correctly and what needs to change. If we continue to run our application and there are no more debugger statements, the debugger will run to the end of the application.

Note: the Chalk npm package cannot style the console in the debugger, but this should be OK as it is a little utility to make our application look a little nice when displaying outputs, therefore, not essential to our application code.

Once we run the code in the debugger to the end of our application and close down the developer tool, we should see that on the <chrome://inspect> page, there are no longer any Remote Target available. If we wanted to run through the program again and debug things, we can run a single command in the terminal called restart:

```
< Waiting for the debugger to disconnect...
debug> restart
```

This is going to restart the program again and we should see a new remote target in the <chrome://inspect> page again that we can click to inspect to run and debug the application again. To shut down the command we can press **ctrl + c** twice on our keyboard within the terminal to exit the process to bring us back to the command line to do whatever we want.

We can fix our code after debugging and remove the debugger statement once we are done fixing our code. These are the couple of tools we can use when debugging our Node.js applications.

Errors Messages

We can explore some of the common error messages we may see when our application runs:

ReferenceError: variableName is not defined

After a new line maybe from 3 to 5 lines down in the terminal when we run our code, this provides the actual error message for why the application failed. This provides the explanation from V8 engine as to why our application code could not run.

A ReferenceError is a whole category of errors, but should follow by a explicit message to what the error relates to. This lets us know why things failed but it is not really letting us know why and so it is up to us as developers to figure out what is going wrong. So in the above, is variableName not defined because we spelt it wrong or is it not defined because the function/method did not give us anything back what we thought it would have and therefore when we referenced it later, the code broke. It is still up to us to really work through the problem that caused the error, but the error could get us to the correct part of the code.

Below the error message, we have a stack trace which contains a trace of all of the functions that are running to get to the point where the code breaks. This will point to the file that contains the function as well as the line and the column number. Therefore the first line after the error message contains the most useful information in the stack trace. As we go down the stack trace we typically get less actionable information.

Further down the list we get towards we can view the where the function was called by another function/method in our application code. After our code, we then get into the internals for example code running from packages such as yargs and the further we get down the stack trace list, the closer we get to the node internals codes functions.

To conclude, when we are looking at the stack trace, we want to start from the top to bottom because the top lines in the terminal contains the most explicit pertinent information related to the error.

We now should know a little about how to work with error messages and how to use the node debugger tools to debug our node.js applications using the Chrome developer tools. These are both techniques we can use to fix our issues faster. Over time as we get more comfortable with Node, we as developers should make fewer issues for ourselves to solve and the ones we do create, we should be able to be solved faster. This like everything else, it is a skillset that needs to be built up over time.

Section 5

ASYNCHRONOUS NODE.JS

Introduction

If we read a few articles of what Node is, we are likely to come across the same four terms recurring over and over again. These four terms are: Asynchronous, Non-Blocking, Single Threaded and Event Driven.

These terms are accurate ways of describing Node.js but the question is what do these terms mean and how are they going to impact the node applications we are building. We will explore these four terms within section 5 of this document.

Asynchronous Basics

Non-blocking means the application can continue to run other tasks while it is waiting for some long running I/O process to complete. This is what makes node.js fast and efficient.

Below is a basic example of asynchronous code to understand how node.js operates asynchronous tasks.

In a synchronous programming model, one line runs after the next sequentially regardless of how long each line takes to execute. For example:

weather-app > app.js:

1. console.log('starting')
2. console.log('stopping')

In the above the first line will execute and print to the console and once completed the next line will execute and print to the console. We can now take this example and write some asynchronous code in-between the two console.log to see the behaviour of the asynchronous model.

weather-app > app.js:

1. console.log('starting')
2. setTimeout(() => { console.log('2 second timer') }, 2000)
3. console.log('stopping')

The setTimeout is the most basic asynchronous code that node.js provides us. The setTimeout is a function which allows us to run some code after a specific amount of time has passed. This function takes in two arguments and both are required. The first is a function and the second is a number which represents the number of milliseconds we want to wait before the callback function gets executed. Therefore, 2000 millisecond is equal to 2 seconds.

In the above example, the synchronous model will execute the first line and print to the console, it will then

move onto the second line and execute the code which will wait two seconds before printing to the console. It will then finally move onto the third line and execute the final code and print to the console. So the output would look something like:

```
:~/weather-app$ node app.js
starting
2 second timer
stopping
```

However, asynchronous model will output the logs in a different order. This is because the application will continue to execute the next code without having to wait for the previous code to complete its execution. We will notice the order of the console.logs will be different:

```
:~/weather-app$ node app.js
starting
stopping
2 second timer
```

To conclude, in a synchronous model we needed to wait two seconds before the programme would continue to the next line of code, while a asynchronous non-blocking model; node.js can continue to do other things i.e execute the next line of code below while it waits for the two seconds to pass. The above setTimeout helps demonstrate the basics of asynchronous programming; however, if we imagine a database request to fetch data for one user, it would be nice to do other things while we are waiting for the

time it takes to get that information back from the database. With a non-blocking model, we are able to do a lot of different tasks at the same time for example waiting for 60 database requests to come back while still issuing more requests.

Below is another example of asynchronous programming, but one which provides a weird result:

weather-app > app.js:

1. `console.log('starting')`
2. `setTimeout(() => { console.log('2 second timer') }, 2000)`
3. `setTimeout(() => { console.log('0 second timer') }, 0)`
4. `console.log('stopping')`

```
:~/.../weather-app$ node app.js
```

```
starting
```

```
stopping
```

```
0 second timer
```

```
2 second timer
```

We would have expected '0 second timer' to print to the console before 'stopping' because we are not waiting for any timer for the function to execute, however, this is clearly not the case. We will explore the internals of Node.js that are responsible for actually asynchronous programming to work, so that we can understand how Node.js is running the code behind the scenes and understand our code executions.

Call Stack, Call Queue and Event Loop

We will explore the internals of Node.js to understand how asynchronous code are executed so that we can answer questions such as why does a `setTimeout` delay not prevent the rest of the programme from running or why does a `setTimeout` of zero milliseconds cause the function to run and execute after code below it as seen in the previous examples. This should hopefully provide us a better visual understanding of what is occurring behind the scenes and exposure to various forms of asynchronous programming.

In the first of three examples, we will see the behind the scenes for a simple synchronous code running in Node:



Even in a simple synchronous code, there is still a lot going on in the background. To break down the above image, we have:

1. In the top left corner we have the node file that is executing. This very simple script that creates a couple of variables and then we print a message to display to the screen.
2. In the bottom left corner we have the terminal output for the script. So when the script prints the message, it will appear in this window.
3. On the right hand side we have all of the internals running behind the scenes in both node and V8 engine. These are the Call stack, Node APIs and the Event Loop, all of which work in tandem to get our asynchronous programmes running.

For a simple synchronous script, the only behind the scenes we are concerned with is the Call Stack.

The Call Stack is a simple data structure provided by the V8 JavaScript engine. The job of the Call Stack is to track the execution of our programme and it does that by keeping track of all of the functions that are currently running.

The data structure for the Call Stack is very simple: we can add something on to the top of the list and we can remove the top item and that is it. It is best to think of the Call Stack as a can of tennis balls. We can add a tennis balls to the top by dropping it into the can and if we want to remove a tennis ball from the can, we would have to remove the top one's first. We cannot remove a tennis ball in the middle without first removing the ones above it and we cannot add one to the bottom if there are already balls inside the can.

We can explore how the Call Stack helps us to run our above script on the left hand side. The first thing that happens with our script, it is wrapped in the main() function that is provided by node.js - this is not a function created by us, but a anonymous function created and provided by node.js. This anonymous function is often referred to as the main function for the programme.

This first thing that happens is that the main() function gets pushed onto the Call Stack dropping all the way to the bottom, since there is nothing else inside the Call Stack. When something is added to the call stack, it means that it will get executed.

The screenshot shows a terminal window with a dark background. On the left, a script is displayed with the following code:

```
1 const x = 1
2
3 const y = x + 2
4
5 console.log('Sum is ' + y)
```

A blue arrow points to the first line of the script. On the right, a sidebar titled "Call Stack" shows a single purple bar labeled "main()".

This main function will kick things off with line 1 and move its way down the script lines. Therefore, this will create a const of x with a value of 1, then a const of y with a value of 3 ($1 + 2 = 3$) and then it will move onto the last line to log 'Sum is 3' within the console/terminal. The log() is a function; whenever we call on a function, the function gets added onto the Call Stack. Therefore, the log() function is now added to the Call Stack to be executed and we now have two functions on the Call Stack, which are main() and log() functions.

The screenshot shows a code editor with a dark theme. On the left, there is a terminal window displaying the output of a script. The script content is:

```
1 const x = 1
2
3 const y = x + 2
4
5 console.log('Sum is ' + y)
```

The terminal output shows:

```
1 Sum is 3
2
3
4
5
6
```

On the right, there is a "Call Stack" section. It contains two purple rectangular boxes stacked vertically. The top box contains the text "log('Sum is 3')". The bottom box contains the text "main()". A blue arrow points from the bottom of the "log" box towards the bottom of the "main" box, indicating the flow of execution.

Once the log function executes, this will print the 'Sum is 3' in the terminal which is what we would see when we run the programme.

When a function finishes by either running to the end or returning a value, the function gets removed from the Call Stack since it is no longer executing.

Therefore, at this point the log function has done its job by printing something to the console. It is now going to get finished by being removed from the Call Stack.

The blue arrow will then go to the next line of the programme which is actually the end of the script, which means the main function is going to finish as well and be removed from the Call Stack leaving us with an empty Call Stack. At this point the programme is now completed.



In the second example, we still have a synchronous script but there is a little more complexity with multiple function calls.



In the script we define a `listLocations` function which takes in an array and for each location, it loops through the item, printing the location to the console. Below we define the location array stored in a variable and pass it into the function.

We can now analyse how node.js will run this programme in the background should we run this script.

The first thing that will execute is the main() function which will be added to the Call Stack, this will allow the script to start executing line by line.

This will define the listLocations variable creating the function which is the variable's value. At this point we are not calling the function so it is not going to appear on the call stack.

```
●●●  
1 const listLocations = (locations) => {  
2   locations.forEach((location) => {  
3     console.log(location)  
4   })  
5 }  
6  
7 const myLocations = ['Philly', 'NYC']  
8  
9 listLocations(myLocations)
```

From there we move onto line 9, which is where we call the listLocations function passing in the myLocations array object. This is indeed a function call and so we are going to see listLocations added onto the Call Stack.

Call Stack

```
●●●  
1 const listLocations = (locations) => {  
2   locations.forEach((location) => {  
3     console.log(location)  
4   })  
5 }  
6  
7 const myLocations = ['Philly', 'NYC']  
8  
9 listLocations(myLocations)
```

main()

Call Stack

```
●●●  
1 const listLocations = (locations) => {  
2   locations.forEach((location) => {  
3     console.log(location)  
4   })  
5 }  
6  
7 const myLocations = ['Philly', 'NYC']  
8  
9 listLocations(myLocations)
```

listLocations([...])

main()



```
1 const listLocations = (locations) => {  
2   locations.forEach((location) => {  
3     console.log(location)  
4   })  
5 }  
6  
7 const myLocations = ['Philly', 'NYC']  
8  
9 listLocations(myLocations)
```

Call Stack

```
forEach(...)  
listLocations([...])  
main()
```

The listLocations function will now start to execute which is defined between line 1 to 5. The only thing inside of this function is a call to the forEach method looping over each location provided. The forEach is a function call, which is also going to be added onto the Call Stack.

Call Stack

```
anonymous('Philly')  
forEach(...)  
listLocations([...])  
main()
```

This function is going to run one time for each location, it gets access to the location and it will print it to the console/terminal. The forEach function is going to be responsible for calling this function multiple times. The first time it gets called, it gets added onto the stack. We would see the anonymous function is called with the argument 'Philly'.



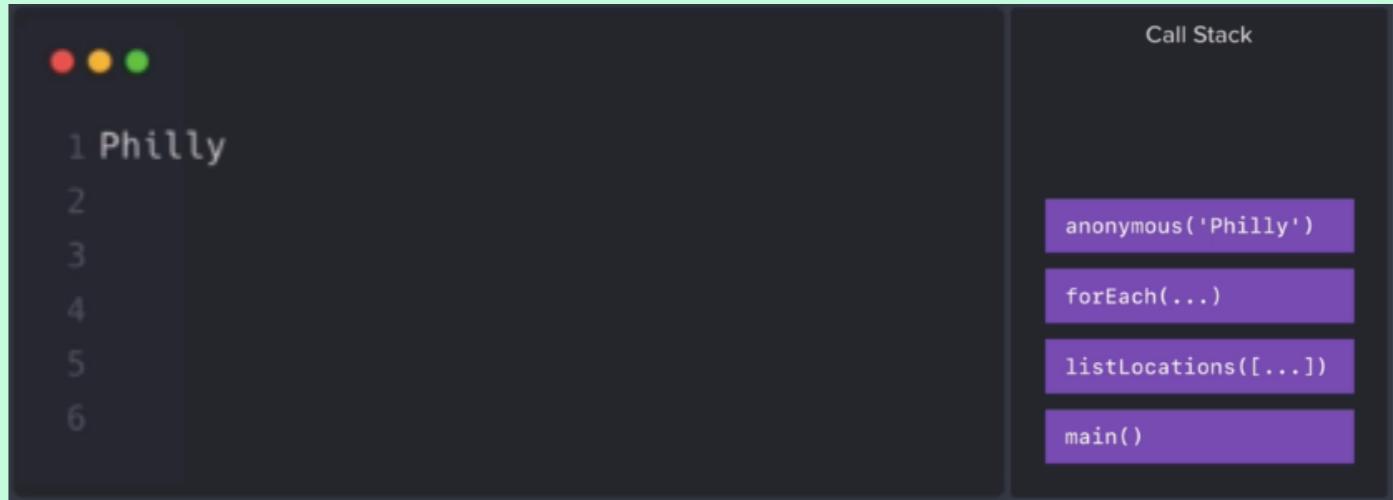
```
1 const listLocations = (locations) => {  
2   locations.forEach((location) => {  
3     console.log(location)  
4   })  
5 }  
6  
7 const myLocations = ['Philly', 'NYC']  
8  
9 listLocations(myLocations)
```

Call Stack

```
console.log('Philly')  
anonymous('Philly')  
forEach(...)  
listLocations([...])  
main()
```

From there we call on the log function to print to the console/terminal which is also then added to the Call Stack.

We will now see Philly printing to the terminal and once this has finished, the `console.log('Philly')` function will now be removed from the Call Stack.

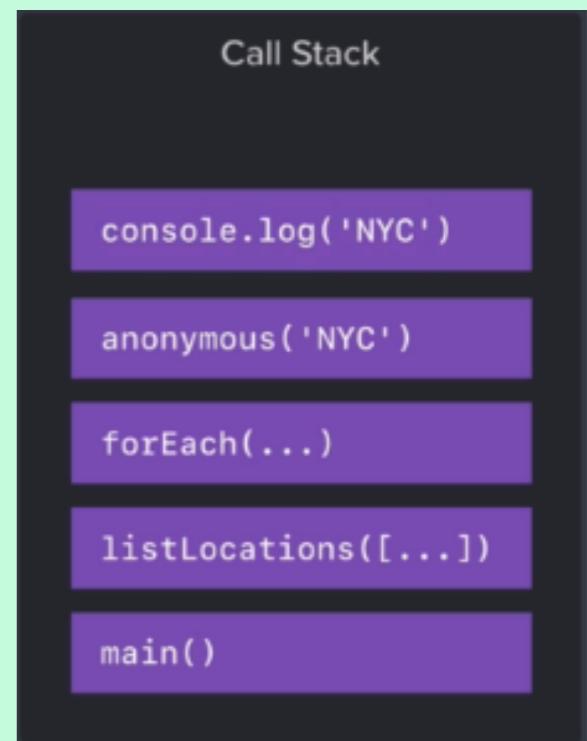


Call Stack

```
forEach(...)
listLocations([...])
main()
```

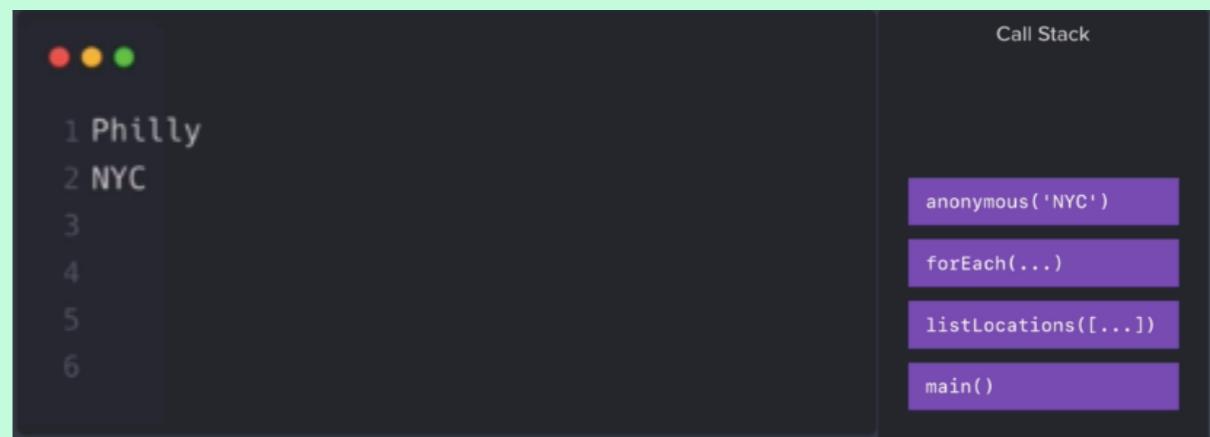
When this log finishes, this is also the end of our anonymous function. This anonymous function will also now pop off the Call Stack as well.

The `forEach` does not come off the Call Stack because it is not completed its execution. It still has to call the anonymous function again to make sure NYC prints to the console. We are going to see a new anonymous function added onto the Call Stack and from there we are going to see another log function added as well.



Once the log executes and prints NYC to the console, the log function will be removed from the Call Stack.

This is also also the end of the anonymous function which will finish and also removed from the Call Stack.



At this point, the `forEach` function has gone through the process two times because there are only two items in the array. The `forEach` is now completed and will also be removed from the Call Stack as well.



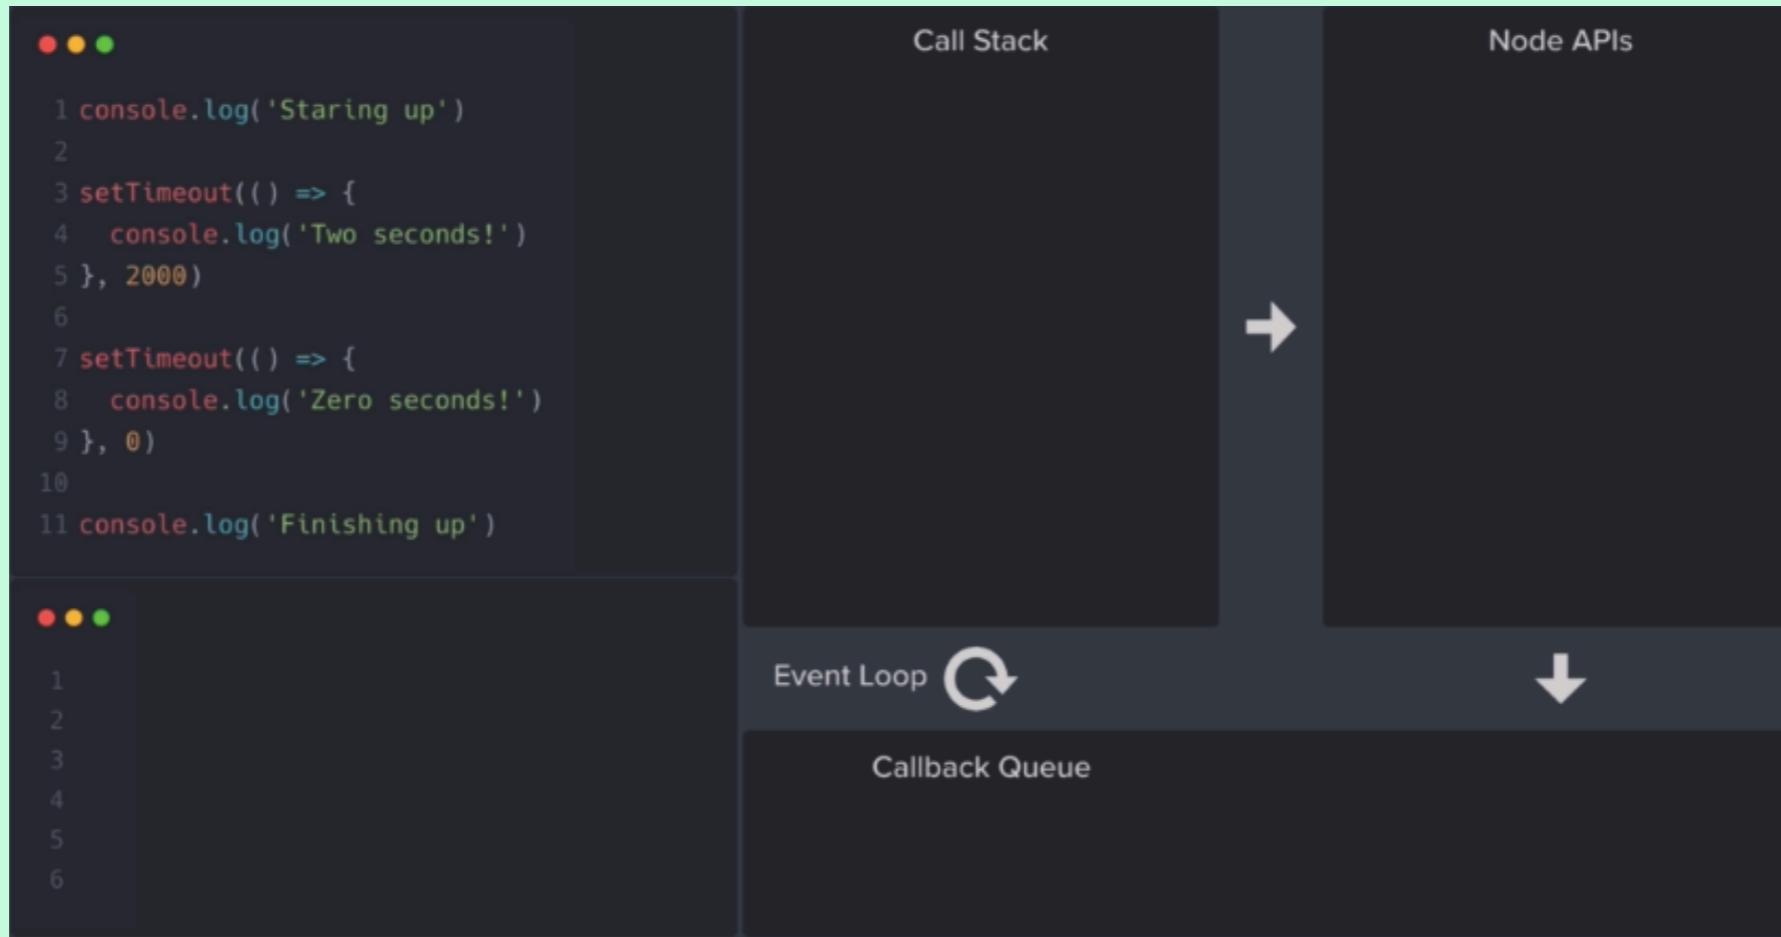
Now that the `forEach` is completed, the `listLocations` function is also completed because it is the only thing the `listLocations` function does i.e. it calls the `forEach` method. The `listLocations` function will also now be removed from the Call Stack.

This now brings us back to line 9 of our script which is the last line of our programme which also means that the `main()` function has now finished and will be removed from the Call Stack. Our Script has completed its execution and the final result is Philly followed by NYC printing to the console/terminal.

A screenshot of a terminal window showing the output of a script. The window has a dark background with three small colored dots (red, yellow, green) in the top-left corner. The output is displayed in white text on the screen:

```
1 Philly
2 NYC
3
4
5
6
```

In the third and final example, we have our asynchronous script we explored at the beginning of this chapter. We are now going to explore how it runs behind the scenes to understand why we were seeing things printed to the console in the order that we were seeing them.



Since the code in this example is asynchronous, we will be using the Call Stack along with the Node APIs, the Callback Queue and the Event Loop all working together to get our application running.

Like with our synchronous examples, the first thing added to the V8 engine's Call Stack is the `main()` function which starts the execution of our programme. On line 1, it calls on the `log` function which also gets added to the Call Stack.

```
1 console.log('Starting up')
2
3 setTimeout(() => {
4   console.log('Two seconds!')
5 }, 2000)
6
7 setTimeout(() => {
8   console.log('Zero seconds!')
9 }, 0)
10
11 console.log('Finishing up')
```

Call Stack

Call Stack

1 Starting up
2
3
4
5
6

console.log('Sta...')
main()

main()

Our message will show up in the console/terminal and the log function is removed off of the Call Stack once it has completed its execution. The next step will move onto line 3 of our script which calls on the setTimeout function which gets pushed onto the Call Stack.

```
1 console.log('Starting up')
2
3 setTimeout(() => {
4   console.log('Two seconds!')
5 }, 2000)
6
7 setTimeout(() => {
8   console.log('Zero seconds!')
9 }, 0)
10
11 console.log('Finishing up')
```

Call Stack

Call Stack

1 Starting up
2
3
4
5
6

setTimeout(..., 2000)
main()

main()

The setTimeout function is not part of the JavaScript programming language and we are not going to find its definition anywhere in the JavaScript specification and V8 has no implementation for it.

Instead, it is Node.js which creates an implementation of setTimeout using C++ and provides it to our Node.js scripts to use.

The setTimeout function is a asynchronous way to wait a specific amount of time and then have a function run. Therefore, when we call setTimeout, it is really registering an event with the Node.js APIs. This is an

event callback pair, where the event in this case is to wait two seconds and the callback is the function to run. Another example of a event callback pair might be to wait for a database request to complete and then run the callback that does something with the data.

Therefore, when we call the setTimeout function, a new event gets registered in the Node APIs which is the setTimeout callback waiting for two seconds. At this point in the process, the two seconds timer starts ticking down. While we are waiting for those two seconds to complete, we can continue to perform other tasks inside of the Call Stack.

It is important to note that JavaScript itself is a single threaded programming language. You can do one thing at a time and the Call Stack enforces that. We can only have one function on the top of the Call Stack which is the task we are executing. There is no way to execute two tasks at the same time. This does not mean Node.js is completely single threaded. The code we run is indeed still single threaded but Node.js uses other threads in C++ behind the scenes to manage our events, which allows us to continue running our application while we are waiting for the event to complete. This is the non-blocking nature of Node. The setTimeout is not blocking the rest of our application from running.

Node APIs

setTimeout (2 sec)



The main() function moves onto line 7 which is another setTimeout function call which gets added to the Call Stack again.

This is going to then register yet another event callback in the Node API area where the callback is waiting zero seconds. We now have two Node APIs waiting in the background. We can continue to do other things while both of the callbacks are waiting for the event to complete (the first case is two seconds and the second case is zero seconds).



The zero seconds are up, therefore the callback needs to get executed. This is where the Callback Queue and the Event Loop come into play. The job of the Callback Queue is simple; its job is to maintain a list of all of the callback functions that are ready to get executed. So when a given event is done, i.e. the zero second timer is completed, that callback function is going to be added onto the Callback Queue. This is a simple line, whereby the function gets added at the end of the line and the front line gets executed first working its way down the list.

Since there are no items in the Callback Queue list, the callback for the zero timer event gets added to the front of the Queue. We now have this callback and it is ready to get executed, but before it can be executed, it needs to be added onto the Call Stack. The Call Stack is where functions go to get run/executed. This is where the Event Loop comes into play.

The Event Loop looks at two things, the Call Stack and the Callback Queue. If the Call Stack is empty it is going to run items from the Callback Queue. So at this point the Event Loop says that it knows that a callback function got added to the Callback Queue, however, the Call Stack is not empty so it cannot execute the Callback function and therefore the reason why the function does not get executed right away.



A screenshot of a debugger interface showing the Call Stack. On the left, there is some code with line numbers 1 through 11. Line 11 is highlighted with a blue arrow. On the right, the 'Call Stack' section shows two frames: 'console.log('Fin...')' at the top and 'main()' below it. The background of the stack frames is purple.

```
1 console.log('Starting up')
2
3 setTimeout(() => {
4   console.log('Two seconds!')
5 }, 2000)
6
7 setTimeout(() => {
8   console.log('Zero seconds!')
9 }, 0)
10
11 console.log('Finishing up')
```

At this point the `main()` function continues to run. The next thing we see is that line 11 of our programme is going to run, which is a call to `log` so the function gets added onto the Call Stack and out message gets printed to the console/terminal.



The log function gets removed from the Call Stack and at this point the `main()` function has completed and will also get removed from the Call Stack.

With our regular synchronous scripts completed, this is when the programme has actually finished. The end of `main()` signifies the end of the application. This is not the case with our asynchronous programmes.

The Event Loop can now start to perform its job. It can see that the Call Stack is empty and there are callback functions within the Callback Queue. It takes that item and moves it up into the Stack so that the callback can run.

At this point it can run the function on line 8 which is a single call to log which gets added to the Call Stack. The message gets printed onto the console/terminal.

```
1 console.log('Starting up')
2
3 setTimeout(() => {
4   console.log('Two seconds!')
5 }, 2000)
6
7 setTimeout(() => {
8   console.log('Zero seconds!')
9 }, 0)
10
11 console.log('Finishing up')
```

Call Stack

```
1 Starting up
2 Finishing up
3 Zero seconds!
4
5
6
```

console.log('Zero...')

Callback (0 sec)

Call Stack

Callback (0 sec)

Once the log completes, the function is removed from the stack and this is also the end of the Callback function which would also get removed from the Call Stack.

This is the reason for why we were seeing Zero seconds! after Finishing up in the terminal. None of our asynchronous callbacks is ever going to run before the main() function is done.

At this point the programme is not done as it has to wait for the two seconds event for the callback to be added to the Callback Queue .

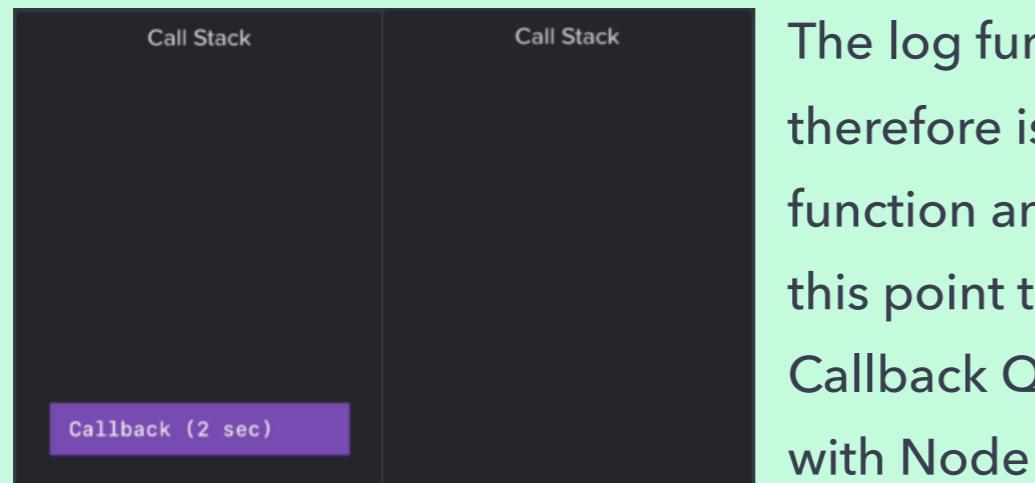


The Event Loop will notice that the Call Stack is empty which means it is ready to run.

The Event Loop takes the callback item adding it to the Call Stack to be executed.



The callback is defined on line 4 which is a call to the log function, which is added to the Call Stack which then gets executed and printed to the console/terminal.



The log function finishes when it prints to the console/terminal and therefore is removed from the call stack which also finishes the callback function and therefore the callback is also removed from the Call Stack. At this point the programme is complete. The Call Stack is empty, the Callback Queue is also empty and there are no other Events registered with Node APIs.

We now know why we got the messages printing in the order we saw them and how Node.js runs our code in the background for both the synchronous and asynchronous model.

Making HTTP Requests

HTTP request allows our application to communicate with the outside world. If we want to get real time weather data into our application, we are going to have to make a HTTP request. If we want to send an email to someone from our application, this is also going to be another HTTP request, the same is true with sending a text message using something like the Twilio API which also requires a HTTP request.

HTTP request is at the core of building real world applications that actually communicate with the outside world. In the above three examples, it is our Node.js application making HTTP request to another companies server to get some task done. This would mean somewhere in our code we are going to specify the URL we want to make a request to, which is provided in the APIs documentation, and we are going to fire that request off sending some data (possibly) and getting a response back. For example, to get a weather information, we would send the location we want for the weather information for and we would get back that weather information which we can use in any way we would like in our application.

To practice making HTTP Requests with some weather data, we can use the website: <https://darksky.net/dev> which is the developer portal for the darksky weather application, allowing us to integrate their data into our application. We are required to signup; however, this API is free to use with 1,000 free requests every day. This is more than enough which will help us understand how to make HTTP requests in our Node.js applications and how we can use the returned data.

Once we signup to the API and login we should see a screen with our API secret key and a sample API Call example as seen in the below example:

The screenshot shows the 'Your Account' section of the Dark Sky API settings. At the top, there are navigation links: 'DOCS', 'FAQ', 'CONSOLE', 'ACCOUNT SETTINGS', and 'LOG OUT'. A success message 'Your secret key has been reset.' is displayed. The 'Your Secret Key' section contains a placeholder text '774f2be3dbdbff72e8bab9cf673a2e70' and a 'RESET SECRET KEY' button. Below this, the 'Plan Type: Trial' section is shown, with a note about the trial limit and a blue 'UPGRADE' button. The 'Sample API Call' section includes a link to a sample API endpoint.

The Sample API Call url link consists of our secret key followed by the longitude and latitude of the location we want to get the data from. If we click on this link to make the HTTP request to the darksky API, we would see the data we would get back from the sample call request.

```
{"latitude":37.8267,"longitude":-122.4233,"timezone":"America/Los_Angeles","currently":{"time":1567348341,"summary":"Clear","icon":"clear-day","nearestStormDistance":239,"nearestStormBearing":163,"precipIntensity":0,"precipProbability":0,"temperature":65.73,"apparentTemperature":65.73,"dewPoint":54.44,"humidity":0.67,"pressure":1013.4,"windSpeed":3.14,"windGust":4.14,"windBearing":252,"cloudCover":0.02,"uvIndex":0,"visibility":9.297,"ozone":294},"minutely":[{"summary":"Clear for the hour.","icon":"clear-day","data":[{"time":1567348320,"precipIntensity":0,"precipProbability":0}, {"time":1567348380,"precipIntensity":0,"precipProbability":0}, {"time":1567348440,"precipIntensity":0,"precipProbability":0}, {"time":1567348500,"precipIntensity":0,"precipProbability":0}, {"time":1567348560,"precipIntensity":0,"precipProbability":0}, {"time":1567348620,"precipIntensity":0,"precipProbability":0}, {"time":1567348680,"precipIntensity":0,"precipProbability":0}, {"time":1567348740,"precipIntensity":0,"precipProbability":0}, {"time":1567348800,"precipIntensity":0,"precipProbability":0}, {"time":1567348860,"precipIntensity":0,"precipProbability":0}, {"time":1567348920,"precipIntensity":0,"precipProbability":0}, {"time":1567348980,"precipIntensity":0,"precipProbability":0}, {"time":1567349040,"precipIntensity":0,"precipProbability":0}, {"time":1567349100,"precipIntensity":0,"precipProbability":0}, {"time":1567349160,"precipIntensity":0,"precipProbability":0}, {"time":1567349220,"precipIntensity":0,"precipProbability":0}, {"time":1567349280,"precipIntensity":0,"precipProbability":0}, {"time":1567349340,"precipIntensity":0,"precipProbability":0}, {"time":1567349400,"precipIntensity":0,"precipProbability":0}, {"time":1567349460,"precipIntensity":0,"precipProbability":0}, {"time":1567349520,"precipIntensity":0,"precipProbability":0}, {"time":1567349580,"precipIntensity":0,"precipProbability":0}, {"time":1567349640,"precipIntensity":0,"precipProbability":0}, {"time":1567349700,"precipIntensity":0,"precipProbability":0}, {"time":1567349760,"precipIntensity":0,"precipProbability":0}, {"time":1567349820,"precipIntensity":0,"precipProbability":0}, {"time":1567349880,"precipIntensity":0,"precipProbability":0}, {"time":1567349940,"precipIntensity":0,"precipProbability":0}, {"time":1567350000,"precipIntensity":0,"precipProbability":0}]]}
```

This is JSON data coming back and this contains all of the weather information for the provided location. When it comes to exchanging data between services, JSON is the standard format for getting this done.

This JSON data is easier for a machine to use and we can use all of the values inside of the data request to provide our users with a detailed forecast of their weather.

We can make the same request as provided in the Sample API Call within our Node.js application. To make a HTTP request, there are a few different things we could do:

1. We can use the core node modules, but this is a very low level (barebones way) which requires us to write a lot of unnecessary code to get everything working together.
2. There are a bunch of npm modules that are essentially wrappers around the core module making it much easier and streamlined to make HTTP requests.

We will use the second option of using the npm module called Request (<https://www.npmjs.com/package/request>). We will later learn how to create HTTP request using the barebones core module approach.

In the terminal we would need to initialise our project as a npm project by running the command in the terminal within the root project directory. The -y flat answers yes to all the setup questions using the default values:

```
:~/.../weather-app$ npm init -y
```

We can then run the following command to install the request npm package:

Once we have our package.json file generated with all of the default values in place, we can now run the following command to install the request npm package. The i flag is short for install:

```
:~/weather-app$ npm i request@2.88.0
```

Once the package has been installed locally in our project, we can go into our app.js file to load and start using the request module to make a HTTP request to our darksky API via our application.

weather-app > app.js:

```
1. const request = require('request')
2. const url = 'https://api.darksky.net/forecast/774f2be3dbdbff72e8bab9cf673a2e70/37.8267,-122.4233'
3. request( { url: url }, (error, response) => {
4.     console.log(response)
5. })
```

This is going to now work because we have the package installed in our package.json file as well as all the necessary modules within our node_modules folder to get the request package to work. Therefore, we have request itself as well as all of its dependencies showing up inside of the node_module folder.

To use request module, we need to work with the same URL we have in the Sample API Call example in the browser where we got the JSON data back. So instead of visiting the URL and viewing the data in the

browser, we are going to visit via the request module and we will get that data back as a variable that we could access in our code.

The request function takes in two arguments, the first is an options object which outlines what we would like to do i.e. this is where we would provide the url and other information and the second argument is a function to run once we actually have that response i.e. once the data is available for us to use in our application.

In the first argument we use the url property and pass in the url which we stored in a variable called url. The second argument is the function to run when we are actually getting the response back which we used an anonymous callback function. This function will run when either the response returns the weather data or something went wrong (e.g. the machine does not have a network connection) and we were not able to get the weather data.

The anonymous function is called with two arguments of error and the response. If there is an error this will appear in the error argument else the error argument will return undefined and the response argument will allow us to actually access the response data. The response data is going to contain everything about the response which is way more information than just the JSON data.

If we run the application in the terminal we should see the successful response logged to the terminal:

```
:~/weather-app$ node app.js
```

We are going to see hundreds of lines of information printed to the terminal and response has many properties we could use and those properties have other properties and so on. The string of data in green is the string JSON data that we would want to parse and access which is contained in the body property.

We already know how to parse JSON data, below is an example:

weather-app > app.js:

```
1. const request = require('request')
2. const url = 'https://api.darksky.net/forecast/774f2be3dbdbff72e8bab9cf673a2e70/37.8267,-122.4233'
3. request( { url: url }, (error, response) => {
4.   const data = JSON.parse(response.body)
5.   console.log(data)
6.   console.log(data.currently)
7. })
```

We have used `JSON.parse` to parse the data that lives on `response.body` property and store it in a variable called `data`. We can now work with that `data` variable and access any of its properties (e.g. `currently` which contains the current forecast information) to log to the console.

To conclude, we have now accessed our first HTTP API from Node.js allowing us to pull in outside information into our applications.

Customising HTTP Requests

We can customise our HTTP requests using the request module so that request can automatically parse the JSON response for us. We can also explore some options in the darksky API allowing us to do things like change the language and change the units. This allows us to truly customise our requests to suit our applications needs so that we can actually print a forecast.

To make the request module automatically parse the JSON response for us, we can do this by customising the options object that we pass as the first argument in our request function. The list of all available options are available in the request documentation. We can add a second property of json and set this value to true.

weather-app > app.js:

```
1. const request = require('request')
2. const url = 'https://api.darksky.net/forecast/774f2be3dbdbff72e8bab9cf673a2e70/37.8267,-122.4233'
3. request( { url: url, json: true }, (error, response) => {
4.     console.log(response.body.currently)
5. })
```

This would mean the `response.body` is already going to be an object and there is no need for us to create a variable to parse the JSON data. We can now log the `response.body.currently` property to the console as we did before in a fewer line of code.

To print a forecast to the user, we need to figure out what data we wish to include and which to exclude. This may seem difficult viewing the data in the terminal and it may be easier to view it in the browser to explore what data we want to include from a JSON response. There are extensions for browsers that allow us to view nicely parsed/formatted JSON data. A useful extension for the Chrome browser is JSON Formatter (<https://chrome.google.com/webstore/detail/json-formatter/bcjindcccaagfpapjjmafapmmgkkhgoa?hl=en>) which can be installed in our browser. Once the extension has been installed, if we were to view the Sample API Call in the browser and refresh the page, we should now be able to see the JSON data nicely parsed/formatted by the extension. We can switch between the Raw and Parsed JSON data to view the returned request.

```
Raw Parsed

{"latitude":37.8267,"longitude":-122.4233,"timezone":"America/Los_Angeles","currently":{"time":1567354309,"summary":"Clear","icon":"clear-day","nearestStormDistance":255,"nearestStormBearing":167,"precipIntensity":0,"precipProbability":0,"temperature":67.36,"apparentTemperature":67.36,"dewPoint":56.08,"humidity":0.67,"pressure":1014.13,"windSpeed":3.23,"windGust":4.32,"windBearing":242,"cloudCover":0.07,"uvIndex":2,"visibility":9.942,"ozone":294.4},"minutely":[{"summary":"Clear for the hour.","icon":"clear-day","data":[{"time":1567354260,"precipIntensity":0,"precipProbability":0}, {"time":1567354320,"precipIntensity":0,"precipProbability":0}, {"time":1567354380,"precipIntensity":0,"precipProbability":0}], "raw": [{"summary": "Clear for the hour.", "icon": "clear-day", "data": [{"time": 1567354260, "precipIntensity": 0, "precipProbability": 0}, {"time": 1567354320, "precipIntensity": 0, "precipProbability": 0}, {"time": 1567354380, "precipIntensity": 0, "precipProbability": 0}]}], "parsed": [{"summary": "Clear for the hour.", "icon": "clear-day", "data": [{"time": 1567354260, "precipIntensity": 0, "precipProbability": 0}, {"time": 1567354320, "precipIntensity": 0, "precipProbability": 0}, {"time": 1567354380, "precipIntensity": 0, "precipProbability": 0}]}]}
```

All of the properties within this data object is accessible to us from the response.body property. From here we can grab any of its properties such as the currently property which is why the response.body.currently works in our code. We can also use response.body.timezone to get the timezone for the current location. So we should now have a better idea of as to how we are accessing the information.

We can click on the DOCS link in the darksky API page to view the documentation of the API to learn how to customise the options available to us. The Forecast Request link will scroll us to the section which provides us the example request. We can see the [key]/[latitude],[longitude] all showing up in the list with the example response below.

The screenshot shows the 'Forecast Request' section of the DarkSky API documentation. On the left, there's a sidebar with links: Overview, Forecast Request (which is highlighted), Time Machine Request, Response Format, FAQ, Data Sources, API Libraries, Privacy Policy, and Terms of Service. The main content area has a title 'Forecast Request' and a URL 'https://api.darksky.net/forecast/[key]/[latitude],[longitude]'. Below the URL, it says: 'A Forecast Request returns the current weather conditions, a minute-by-minute forecast for the next hour (where available), an hour-by-hour forecast for the next 48 hours, and a day-by-day forecast for the next week.' Underneath that is a section titled 'Example Request' with a code block:

```
GET https://api.darksky.net/forecast/0123456789abcdef9876543210fedcba/42.3601,-71.0589
```

```
{
  "latitude": 42.3601,
  "longitude": -71.0589,
  "timezone": "America/New_York".
```

If we scroll past that example request, we get to the Request Parameters where we can start to customise

exactly what is happening. The key, latitude and longitude parameters are provided directly into the url and are required by the API. All of the optional parameters such as units, lang and excludes are all setup as query strings. This would mean at the end of our url, we are going to setup a query string which starts with a question mark followed by key=value pairs. We can add multiple key=value pairs by sing the & sign. We would use this format to setup the other options we want to customise our HTTP request. It is important that the options comes at the end of the url i.e. after the longitude value.

Below is an example of setting up the unit and language options.

weather-app > app.js:

```
1. const request = require('request')
2. const url = 'https://api.darksky.net/forecast/774f2be3dbdbff72e8bab9cf673a2e70/37.8267,-122.4233?units=si&lang=en'
3. request( { url: url, json: true }, (error, response) => {
4.   console.log(response.body.currently)
5. })
```

We now know how to customise the response we are getting back from the API by adding the query strings to our URL. It is very important to read the API documentation as the setup and customisation of APIs will vary from API to API and therefore a good documentation on how to use the API is crucial.

CHALLENGE:

Print a small forecast to the User

1. Print: "It is currently 58.55 degrees out. There is a 0% chance of rain."
2. Use the temperature and precipProbability properties from the response data to dynamically replace the forecast values for the location in our printed message.
3. Test your work!

SOLUTION:

weather-app > app.js:

```
1. const request = require('request')
2. const url = 'https://api.darksky.net/forecast/774f2be3dbdbff72e8bab9cf673a2e70/37.8267,-122.4233'
3. request( { url: url, json: true }, (error, response) => {
4.     console.log('It is currently ' + response.body.currently.temperature + ' degrees out. There is a ' +
       response.body.currently.precipProbability + ' chance of rain.')
7. })
```

:~/weather-app\$ node app.js

It is currently 68.29 degrees out. There is a 0% chance of rain.

Handling Errors

We can always hope things will go right but that is not always the case. There are plenty of things that could go wrong with our http requests. We are now going to focus on those errors to give the user a better experience when things don't go as expected.

The obvious way of things going wrong when making a http request is if there is no network connection. We can test this by temporarily turning off our computer's wifi connection or unplug the ethernet cable and run a http request script. We would get back a really long error message in the terminal but we would notice a `TypeError: Cannot read property 'body' of undefined`.

In this case when things go wrong and there are no network request the error argument contains a value and response does not. We need to add a little conditional logic to check that there are no error before we try to interact with the response data. The response is undefined and what is causing the error message of body undefined problem.

We cannot fix the lack of network connectivity, the purpose of the error argument is to print out a regular human readable way that the user does not currently have access to the whichever API they are trying to access.

weather-app > app.js:

1. ...
2. `request({ url: url, json: true }, (error, response) => {`

```
3.     if (error) { console.log('Unable to connect to weather service') }  
4.     else { console.log(response.body.currently) }  
5.   })
```

We would add a if statement within our http request function to test whether we have a error returned back. If we console.log(error) this will return the error in a format that is not so bad as to what we had before; however, it still does not contain information that would be useful to someone unless they really knew a good deal about programming. We are going to check for the existence of the error object using the if statement.

If there is an error message we would print a error statement using the console.log so that the user would actually understand such as unable to connect to the API. The else part of the if statement will print the forecast as this would be the block of code that runs if there was no error message meaning we had a successful connection to the API. We now have a little bit of error handling within our application making sure that when things don't go as expected, we at least give the user information that is useful about the error.

In summary, we have these two arguments of error and response. Only one of the argument is ever going to be resolved/populated while the other is undefined. Therefore, if we have a value for error, there is not going to be a value for response and vice versa. This is not the only error we would want to guard against as the other error relates to the users input.

Using the darksky weather API, if we were to take the URL and remove the longitude coordinate to simulate a user related error and paste this URL in a browser. In this case we would get a response with two properties back from the API. We have the code property and the error property on the response body.

In this situation we would need to add a little defensive programming to make sure we are printing a useful error message to the user in this error situation.

It is important to note, where we have an invalid input, the server might respond with data that might be considered an error. This will not populate the error argument object. The error argument is used for lower level OS related errors such as a complete lack of a network connection. If the server we are contacting cannot find the data we need, it is typically come back on the response body argument object. We would need to add an else if clause to our if statement to look at the response and see if there is an error code setup. If there is a error code setup, we can print a different error message, else we would continue to print the forecast data.

weather-app > app.js:

```
1. ...
2. request( { url: url, json: true }, (error, response) => {
3.     if (error) { console.log('Unable to connect to weather service') }
4.     else if (response.body.error) { console.log('Unable to find location') }
5.     else { console.log(response.body.currently) }
6. })
```

The else if statement is checking to see if the response.body has an error property and if it does then the condition is true and will run the code block for the true statement i.e. the message to print the error to the user, else there is no error property and the forecast will be printed instead.

To conclude, there are two types of errors we need to handle and these are low level OS errors and invalid input errors. Using the if statements and the error argument will help us to handle the different edge cases and provide useful error messages to our users.

When working with API, it is important to test the different edge cases for input errors as this could return different results depending on the API setup and the application logic will need to be setup correctly to handle the errors.

The Callback Function

We are going to dive deep into the callback pattern to see how it works and how we can take advantage of it by creating our own. This is important to understand because it is at the core of asynchronous Node.js development. It is worth investing time to learn callback functions so that we are more comfortable with it as we continue to use them.

playground > app.js:

1. `setTimeout(() => { console.log('Two seconds are up') }, 2000)`

In this simple setTimeout example, there is a single callback function and that is the anonymous function we defined and passed in as the first argument to the setTimeout function.

A callback function is nothing more than a function we provide as an argument to another function with the intention of having it called later on. Therefore, in the above example, we are providing a function as an argument to setTimeout with the intention that setTimeout is going to call the function at some point in the future. That's all a callback function is.

In the above we are using the callback function in an asynchronous way. The setTimeout is a node provided API that is indeed asynchronous. This does not mean every time we use the callback pattern, it is actually asynchronous. This is a point of confusion for a lot of people.

We have already proven this; we know that setTimeout uses the callback pattern and it is asynchronous and we know that our array methods such as forEach and filter use the callback pattern but they are synchronous.

playground > app.js:

1. const names = ['Andy', 'Beth', 'Cathy']
2. const shortNames = names.filter((name) => { return name.length <=4 })

In this example the filter method is going to be run each time for each array object; however, we are

returning the names that are 4 or less characters long in the new shortNames array. In this case we are indeed using the callback pattern as well but there is nothing asynchronous about filter. It is not interacting with a native Node API as it is a standard JavaScript code that is completely synchronous.

So far we have only worked with callback functions that we have never defined such as setTimeout and filter.

Imagine in our programme there are four different places where we want to be able to take a location and get the coordinates back using the geocode API. If we have four more location we need to grab the location, we would have to write the code and place it in four or more different places which is not ideal. What we would do instead is call a function called geocode or something alike and all of the code will live in that named function we created. To get this done, we would need to be familiar with the callback pattern, below is how the code would look like using the callback pattern:

playground > app.js:

```
1. const geocode = (address, callback) => {  
2.     const data = { latitude: 0, longitude: 0 }  
3. }  
4.  
5. geocode('London')
```

The above is a function we defined. We can setup this function to take in a callback function as an argument just like in setTimeout and filter, and we can set as many arguments as we would like.

If we pretend that in the future we would get the data back from request, what would we do with it? The goal is to give it back to the caller of geocode. So if we call geocode and pass in an address, we want to get access to the coordinates. There are two ways in which we could get this done:

The first could be a return value from the geocode function as we have seen previously in every function so far. This would mean that we do not need the callback argument:

playground > app.js:

```
1. const geocode = (address) => {  
2.     const data = { latitude: 0, longitude: 0 }  
3.     return data  
4. }  
5. const data = geocode('London')  
6. console.log(data)
```

```
:~/.../playground$ node app.js  
{ latitude: 0, longitude: 0 }
```

The problem with the above is that there is nothing asynchronous happening with our code. Later on we would use request which we know is asynchronous.

We will notice with the below code we would have a unexpected result.

playground > app.js:

```
1. const geocode = (address) => {
2.     setTimeout( () => {
3.         const data = { latitude: 0, longitude: 0 }
4.         return data
5.     }, 2000)
6. }
7. const data = geocode('London')
8. console.log(data)
```

```
:~/playground$ node app.js
```

```
Undefined
```

This will return undefined in the terminal because geocode is not returning anything. In this function there is a single call to setTimeout and there is no return statement directly inside geocode function. Therefore, geocode finishes almost immediately and if you do not return something from a function, we know that JavaScript will implicitly return undefined and so that is what we end up seeing. There is a return in geocode; however, it is nested in another function. That return statement is returning a value from the setTimeout function and not from the geocode function, which is why we are not getting a value back. The return pattern is no longer going to work for us when we start to do asynchronous things inside of our functions. This is where the callback pattern is going to come into play.

The alternative is to provide a callback and get the data from there to get the function back to a working state. We would need to pass in callback as the second argument.

playground > app.js:

```
1. const geocode = (address, callback) => {
2.     setTimeout( () => {
3.         const data = { latitude: 0, logitude: 0 }
4.         callback(data)
5.     }, 2000)
6. }
7. geocode('London', (data) => { console.log(data) })
```

```
:~/playground$ node 4-callbacks.js
{ latitude: 0, logitude: 0 }
```

We know with asynchronous behind the scenes, none of our callback functions i.e. the ones that go to the Node APIs and then come down to the Callback Queue, they do not get executed until the call stack is empty. That means the main() function has to finish. We have to adjust the call to the function call as geocode is not going to have a direct return value. What we do instead is add the callback function as the second argument. The callback function is going to be called at some point in time, the question is what

things are we going to have access to which is unto us to choose. Instead of returning the data in our setTimeout function, we are going to call the callback function passing in the data as the first argument. Therefore, when we call on the geocode function and pass the second argument of the callback function, we are also calling on data as the first argument (the name of the argument in the geocode callback function does not need to match the argument name in the setTimeout callback function as it is only the position that matters). We can now console.log the data.

We are now using the callback pattern in our fictitious asynchronous defined function example.

Note: we could have called the callback argument within const add anything we wanted to as we are not limited to calling it callback.

To summarise, if our function is completely synchronous we are able to use return to get a value out of the function and back to the part of the code that called that function. When our function starts to do something asynchronous, the return is no longer an option. Instead of returning a value, we take a callback in as an argument to our function and we call the callback with the value we want to send back when we have it. This achieves the same goal as the return statement does as it still gets the value back to the code that needs it.

CHALLENGE:

Mess around with the callback pattern

1. Define an add function that accepts the correct arguments
2. Use setTimeout to simulate 2 second delay
3. After the 2 seconds are up, call the callback function with the sum
4. Test your work!

```
add(1, 4, (sum) => { console.log(sum) })
```

SOLUTION:

playground > app.js:

1. const add = (a, b, callback) => {
2. setTimeout(() => {
3. callback(a + b)
4. }, 2000)
5. }
6. add(1, 4, (sum) => { console.log(sum) })

```
:~/playground$ node 4-callbacks.js
```

```
5
```

Callback Abstraction

Building on the previous callback pattern we are going to learn how to create a http request inside of a function. This will allow us to do two very important things:

1. It is going to make it easy to create code that is reusable and easy to maintain i.e. we can call on the function as many times as we need without copying it over and over again.
2. It is going to make doing one thing before or after something else e.g. we can geocode the address and output latitude and longitude to use that to fetch the weather data.

weather-app > utils > geocode.js:

```
1. const request = require('request')
2. const geocode = (address, callback) => {
3.   const url = 'https://api.mapbox.com/geocoding/v5/mapbox.places/' +
    encodeURIComponent(address) + '.json?access_token=pk.eyJ1IjoieXR1YmUtdHV0b3JpYWxzliwi
    YSI6ImNrMDFhaHNiNTBvYnczbXF0bm15aXVxNnQifQ.F9o3SEHg8NI0-qQS5Gt8ng&limit=1'
4.
5.   request({ url: url, json: true }, (error, response) => {
        if (error) { callback('Unable to connect to the services', undefined) }
        else if (response.body.features.length === 0) { callback('Unable to find location...', undefined) }
        else { undefined, { latitude: response.body.features[0].center[1], ... } }
      })
7. }
8. module.exports = geocode
```

We would create a function and define the parameters/arguments required as input for the function to operate. We would require a callback function as an argument because http request is an asynchronous task and we would call on this callback function once we have the latitude and longitude data from the request. When we call on the geocode function we would pass in the address input but also define a callback function where we have access to the results after the geocode process is complete.

When working with callback functions and the callback pattern, it is typical to see two arguments passed to callbacks which is error if things went wrong and data if things went well. This is not enforced but is a common convention that we would see across many different tools, utilities and libraries (as we have already seen with the request npm package). Only one of the two arguments are ever going to have a value while the other is set to undefined, allowing us to conditionally check what is going on.

We have the http request url contained within a variable. This url is going to be static except with the one piece where we would need to add the address we wish to return the latitude and longitude. We can concatenate the string beginning and end and place in between the address input to make the url string dynamic. The encodeURIComponent is a function which is used to concatenate a string but this helps in scenarios where a user searches for a location that contains special characters that actually mean something in a URL structure e.g. ? which will be converted into to encoded version i.e. %3F. This will allow to handle the http request correctly without crashing.

The goal is to create a function that is highly reusable, so when someone calls geocode, they may want to do something different with the error message that log it to the console (e.g. save it to a log file on the file system or email it to a system admin). Instead of logging out the error we can pass it back to the callback and the callback function can choose what to do with the error message. Therefore, we can call geocode() five different times and do five different things with the error message.

In the if statement error condition we used the callback function to pass in the message for the error value and undefined for the data value. We do not need to explicitly specify undefined as JavaScript will define this implicitly if we do not have any values for an argument. In the above example we explicitly defined data to be undefined in the callback function which will also work. These values will be returned to the callback function if the http request was to fail and we can use the callback function to do whatever we want with the returned values. This is callback abstraction where we can abstract the values from the callback function to make our functions more flexible.

weather-app > app.js:

```
1. const geocode = require('./utils/geocode')  
2.  
3. geocode('London', (error, data) => {  
4.     console.log('Error': error)  
5.     console.log('Data': data)  
6. }
```

Callback Chaining

The callback chaining pattern allows us to one function and then then next which is particularly useful when chaining multiple different asynchronous I/O operations/functions for example taking the coordinates returned from one function and using it as the input for the next function.

weather-app > app.js:

```
1. const geocode = require('./utils/geocode')
2.
3. const location = process.argv[2]
4.
5. if (!location) {
6.   console.log('Please provide a location.')
7. } else {
8.   geocode(location, (error, data) => {
9.     if (error) { return console.log('Error': error) }
10.    forecast(data.latitude, data.longitude, (error, forecastData) => {
11.      if (error) { return console.log('Error': error) }
12.      console.log('Data': forecastData)
13.    })
14.  }
```

We have the second function nested in the first function. We start the asynchronous call to geocode and when it is done, the event loop is going to make sure that the callback gets called. From there we are going to start another asynchronous I/O operation called forecast and then wait for that callback to finish and inside of this callback we have access to the final data. This is an example of callback chaining whereby we are chaining together multiple callbacks to do multiple things in a specific order.

We are taking the success data returned from the geocode callback and populating the input argument for the forecast function's latitude and longitude values i.e. the input for forecast comes from the output of geocode. When the forecast resolves it will either print the error or success in the console/terminal from its return error or data argument values.

We use the if statement to check for errors before running the second nested function (in the else clause) to prevent it from running should there be an error in the first function. Alternatively, as demonstrated in the above example, we can use the return statement instead of else to return the error message which will stop the second function from executing after printing the error message to the console/terminal, which is another very common pattern.

The output argument name for data in both geocode and forecast are exactly the same which is a problem because the forecast data argument will overshadow the geocode data argument. To address this issue, all we need to do, is come up with a unique name for one or both of the arguments.

ES6 Object Property Shorthand and Destructuring

We are going to explore some ES6 features that make it easier to work with objects i.e. creating objects and accessing properties on an object.

The ES6 object property shorthand allows us to add values onto an object with a shorthand syntax under certain conditions. Below is the regular ES5 syntax for creating an object and printing it to the terminal:

playground > app.js:

1. const name = 'Alex'
2. const userAge = 20
3. const user = { name: name, age: age, location: 'London' }
4. console.log(user)

The object property syntax can be used when defining an object as seen above. It comes into play when we are setting up a property of which its value comes from a variable of the same name as seen above.

Therefore we can use the shorthand syntax. We remove the colon and the variable leaving what looks like only the property name. This is actually a object property shorthand syntax.

playground > app.js:

1. const name = 'Alex'
2. const userAge = 20
3. const user = { name, age: userAge, location: 'London' }
4. console.log(user)

The property name and variable name must match in order to take advantage of the ES6 object property shorthand syntax else we would end up with a `ReferenceError` message in the console if variable is not defined.

ES6 Destructuring is useful for when we have an object and we are trying to access properties from it. The goal of destructuring is to extract object properties and their values into individual variables.

playground > app.js:

```
1. const product = { label: 'Red notebook', price: 3, stock: 201, salePrice: 'undefined' }
```

In the above example we can destructure the `product` object so that instead of a `product.price` property, we can have a `price` variable with the value of 3, etc. This is useful for when working with complex objects that have a lot of properties we are constantly referencing.

It is nice to have standalone variables that we could easily use, such as `product` and `stock` as seen below. We can access the `label` and `stock` variables throughout the rest of the programme without having to grab it off the object each and every time we want to use it.

playground > app.js:

```
1. const product = { label: 'Red notebook', price: 3, stock: 201, salePrice: 'undefined' }
2. const label = product.label
3. const stock = product.stock
```

The problem with the above example, is that we end up writing a lot of code for each property we want to extract from the object. Therefore, we are going to have multiple lines to pull off multiple values. With ES6 destructuring, we get an improved syntax as seen below:

playground > app.js:

1. const product = { label: 'Red notebook', price: 3, stock: 201, salePrice: 'undefined' }
2. const { label: productLabel, stock, rating = 5 } = product
3. console.log(productLabel, stock, rating)

This syntax seems strange at first, but the more we use it the more easier it becomes to understand. To setup destructuring we create a variable and use the curly brackets followed by the equal sign followed by the name of the object we are destructuring. To create a label and stock variable without accessing it through via `product.label` and `product.stock`, inside of the curly brackets we list out all of the properties we are trying to extract. This creates individual label and stock variables which we can use.

The destructuring syntax makes it easy to extract properties off of an object, creating individual variables that stores the property values. When destructuring we can grab as many properties as we want by listing them out in a comma separated list within the curly brackets and this includes properties that do not exist on the object at all such as rating which value will be undefined by default.

Another nice feature of destructuring is that we can rename the variable we end up creating by following

the property we are extracting with a colon and the new property name. So the new productLabel variable gets its value from the object (which we defined after the equal sign) property (which we defined on the left hand side of the colon).

The last feature we have with destructuring is the ability to setup a default value for the variable should the object not have that property. We do this by adding the equal sign followed by the value after the property as seen in the above for rating. This will override the default undefined for rating as we have now assigned a new default value if the property does not exist on an object. If the object does have the property it will use the value from the object overriding the default we setup should the property not exist on the object.

Finally, we can use destructuring when working with function arguments.

In the below example, we can destructure from our function argument in order to create a local variable that we can use:

playground > app.js:

1. const product = { label: 'Red notebook', price: 3, stock: 201, salePrice: 'undefined' }
2. const transaction = (type, myProduct) => {
3. const { label } = myProduct
4. console.log(label)
5. }
6. transaction('order', product)

We can take this a step further and destructure the argument within the argument list. Therefore, instead of giving a name for the argument, we simply setup the curly brackets and we can destructure whatever we want off of that variable without ever having access to the whole thing. We have access only to the values we choose to destructure.

playground > app.js:

1. const product = { label: 'Red notebook', price: 3, stock: 201, salePrice: 'undefined' }
2. const transaction = (type, { label, stock }) => {
3. console.log(type, label, stock)
4. }
5. transaction('order', product)

When we are using destructuring, we can destructure a standard object outside of the arguments for a function and if we know an argument is an object, we can destructure it right inline as both syntax are valid.

We should now be more familiar with the ES6 object property shorthand and destructuring features which we can now utilise in our applications and to become better developers.

HTTP Requests Without a Library

NPM libraries typically make it easier to do things in node as we have seen with both the request and yargs packages. We do not need those packages to work with command line arguments or to make HTTP requests, but it does make the process a whole lot easier. In this section we are going to explore how to make HTTP request using the modules that node provides.

There are two core modules called HTTP and HTTPS. The documentations for the modules can be found on the node.js website: <https://nodejs.org/dist/latest-v12.x/docs/api/>

There are separate modules depending on the protocol we are using to make our requests. With the npm request package the distinction is not necessary as it abstracts off of that behind the scenes allowing us to easily switch out the protocol we're using without needing to load in a completely separate library.

We can use both these module libraries to create a new server as well as make requests to an existing server. The method we are interested in is the `request` method which appears in both modules and there are two variants (i.e. `.request(options[, callback])` and `.request(url[, options][, callback])`). We use the HTTP module for standard request while we use HTTPS if we are making a request to a secure server.

We will look at the `https` module as an example, since our API requests use the `https` protocol.

playground > app.js:

```
1. const https = require('https')
2. const url = 'https://api.darksky.net/forecast/774f2be3dbdbff72e8bab9cf673a2e70/40,-75?'
3. const request = http.request(url, (response) => {
4.     let data =
5.     response.on('data', (chunk) => {
6.         //console.log(chunk)
7.         data = data + chunk.toString( )
8.     })
9.     response.on('end', () => {
10.         //console.log(data)
11.         const body = JSON.parse(data)
12.         console.log(body)
13.     })
14. })
15. request.on('error', (error) => {
16.     console.log('An error: ', error)
17. })
18. request.end()
```

Since this is a core module, we do not need to install the library and we can simply require the module to load it in. We can use the `request` function to fire off a request. The first thing we need is a URL to make a

request to. We can then make a request using that URL within the `https.request()` function passing in two arguments, first the UR and second a callback function which is called with the response as its argument

The callback we use here is very different from the callback we are use to with the npm request package. Core node modules typically operate at a lower level while npm modules abstract away a lot of that complexity. Since we are using the core node module, we are going to have to setup things that we might not think are necessary. For example, in this callback we do not have access to the complete response body and instead we can go ahead and grab the individual chunks that come through because http data could be streamed in multiple parts. What does this mean for us? It means that we have to listen to each individual chunks to come in and we also have to listen for when all chunks have arrived in the request is done.

We start the process by using `response.on()` which is a function and it allows us to register a handler. There are a few different events we can register callback functions for and one of them is data. We provide the event name as the first argument represented as a string. The second argument we can provide the callback which will fire when new data comes in and we get access to that data via the first and only argument commonly called chunk. This is a chunk of the response which might be the entire thing or it might not depending on how exactly the server has been setup. The other thing we need to setup is to figure out when we are done which we can do via `response.on()` function but this time on the 'end' event handler. When things are over, this callback function is going to run and it does not get back any

arguments because instead by running we just know that we are done. We now need to put both these functions together to figure out how to get the entire response body, parse it from JSON to a javascript object and then actually use it.

The data event handler callback is going to fire when data comes in which could happen once or multiple times; either way we need to take these chunks and add it somewhere we can store it until we have all of the chunks in order to parse it as JSON. We can do this by creating a variable which we called data which we set to an empty string. We use let instead of const as we are going to reassign the value over time changing the string.

If we try to run the code, things are going to hang i.e. it will not do anything and will not return us to the terminal to be able to do any new commands. It is stuck because we do not have a complete request. There is another method we need to use to say we are actually ready to send this off. To do this we need to make a slight change to the programme. What we get back from the https.request method is what we could refer to as the request itself - we store this in the const request variable. This is the variable that comes back from the return value from request. On here there is a method we can use to kick things off which is request.end(). This is another example of how the low level API might not give us the exact tools we were hoping for and it could get a little confusing event with a simple example as there is a lot going on with very little progress.

If we were to `console.log` the chunks we would see either one or a few `console.log` calls printing various Buffers to the terminal. Therefore, the chunk data that comes back is indeed a buffer and we want a string. We would want to convert the buffer to the string using the `toString()` method and add the data on the data variable. We take the existing data value and add the new chunk that is converted from a buffer to a string value. We would end up with access to the entire body response within the data variable.

On the end event handler we are going to have access to the body response data. The data event handler is going to run multiple times, one for each chunk, while end event handler is going to run a single time once things are done. The end data is going to be our long JSON data. From here we could go ahead and parse the data using `JSON.parse` to get an object that we can actually pull properties off of.

This is what it takes to make a request with the core HTTP and HTTPS module and we are not really done yet because we do not have any error handling setup. To setup a error handler, just before `request` and `request.end()`, we would need to setup another `request.on()` listener to listen for the error event handler. Our callback function is going to allow us to do something with the error when this event is triggered. The error is the first and only argument to the callback.

We now have a very basic HTTPS request with the core node module. As we can see this provides everything we need but at a much lower level than we have probably expected. This is why in the real world people are not making requests with the core module and instead they are using libraries like

request, axis or others to make the request process easier. We may be thinking why doesn't Node just change how those core modules work to make it easier to use and look a little more like a library such as a request. The core Node module are supposed to provide the low level implementation and Node comes bundled with NPM because we are suppose to be using NPM modules when building out our applications.

To conclude we took a dive into how we can use the core Node modules to make requests and we have learned a little bit more about why NPM libraries such as requests are so valuable to the Node ecosystem.

Section 6

WEB SERVERS

Introduction

To this point, all of the application code we have seen examples of are for applications accessible via the command line. This is a great way to get started learning Node; however, it is not realistic for users.

We are now going to explore ways for users to interact with our applications via web browsers. We will start to learn one of the most popular NPM libraries called Express which is a web server framework. Express makes it really easy to create web servers with Node.js. Now that we have an understand of the fundamentals of Node.js, this is going to make it easier for us to differentiate between what is Node.js and what is Express.

Hello Express!

We will learn how to build our own Node.js server which is going to allow users an entirely new way to interact with our applications. Instead of needing to run commands from the terminal to interact with our app,

users will be able to visit a URL in the browser and interact with our Node.js application. With the Node.js server, we can serve whatever our application needs such as HTML, CSS and client side JavaScript files, etc. Alternatively, instead of serving up a website we could serve up a HTTP JSON based API similar to the darksky API where we are exchanging JSON data back and forth with the server.

Express is a very popular tool (Node.js server library) used to create Node.js servers. We can find the Express documentation on their website: <https://expressjs.com/>

Express is one of the original npm packages and helped put node.js on the map because it made it so easy to create web servers whether static website or a complex HTTP JSON based API to serve as a backend to serve up a mobile or web application.

In this section we are going to explore how to setup a basic web server whereby we can navigate to a URL in the browser and view an asset that the server has served up.

Inside of the terminal within our project directory, we can run the following code to install the Express npm package within our project directory:

```
:~/.../web-server$ npm i express@4.17.1
```

Once we have successfully installed the Express npm package, we now have a library to create a Node.js script that will configure and start up the server. We can create a sub-directory in our project file called src (short for source) and this will be the directory where we would end up putting all of our Node.js scripts. This will help keep things more organised rather than keeping all our files in the root directory. This will help us create a much more complex directory structure which will allow us to scale a little better as the application starts to grow. As things become more complex, it is important to stay more organised.

Within the src folder we will create a file called app.js which will act as the starting point to our node application. We will use this file to load in express, configure it to serve something up and then start the server.

web-server > src > app.js:

```
1. const express = require('express')
2. const app = express()
3. app.get('', (req, res) => {
4.   res.send('Hello Express!')
5. })
6. app.listen(3000, () => {
7.   console.log('Server is up on port 3000')
8. })
```

We first load in the express package by setting a const variable and requiring the express library exports. The express library exposes a single function. Express is actually a function as opposed to something like an object and we call it to create a new express application.

We create a variable, for example app, to store our express application. To generate the application we simply call on the express function. The express function does not take in any arguments, instead we configure our server by configuring various methods provided on the application itself. This mean we are all done setting up a express application and we can start to tell our express application what exactly it should do.

If we imagine that we own the following domain called app.com for our application. When someone visits this URL we want to show the user something such as a homepage for the company website. We will also have other pages such as app.com/help or app.com/about etc. Therefore, we have a single domain and all of it is going to run on a single express server. What we have setup are multiple routes such as the root, / help and /about routes and we can add many other routes to our application.

We would configure our server to send a response when someone tries to get something from a specific route. We set this by setting up a method on the express application object which we called app. This method is called .get() which allows us to configure what the server should do when someone tries to get the resource at a specific URL. The .get() method takes in two arguments, the first is the route and the second is the function. The route is represented by a empty string. The function is where we describe what

we want the server to do when someone visits the particular route. The anonymous function gets called with two very important arguments, the first is an object containing information about the incoming request to the server which is common called req which is short for request. The second argument is the response common called res which is short for response. The response contains a bunch of methods allowing us to customise what we are going to send back to the requestor.

Within the code block we can go ahead and look at the request to figure out what we want to do such as read data from the database, create some html, etc. We can use various methods on response to actually send a response back. The above example provides a very basic text response displaying some text in the browser using the .send() method.

The .send() method allows us to send something back to the requestor. Therefore, if the requestor is making a request from code using something like the npm request library, they will get the response text back and if they are making a request from the browser, this is the text response that is going to display in the browser window. In the above we now have something that will show up at the root URL i.e. the text Hello express! displayed in the browser window.

The last thing we need to do is actually start the server up. To do this we need to use a method on app which we will ever use a single time in our application. This method is called .listen() which will start up the server and have it listen on a specific port. A common development port is port 3000.

It is important to note that port 3000 is obviously not the default port. When we visit a website, we do not provide the port because there are default ports for example, for a HTTP based website the default port is 80. As we explore production deployment in later sections, we will learn how to use those ports. In local development environment where we are viewing things on our local machine, port 3000 is going to work really well.

The second argument we pass to the `.listen()` method is a callback function which just runs when the server is up and running. The process of starting up a server is a asynchronous process, although it will happen almost instantly. We can print a message to the console to let the person who is running the application know that the server did start correctly. This message will not display in the browser.

Now that we have our application server loaded and configured, we can start up the server by running the following code in the terminal:

```
:~/.../web-server$ node src/app.js  
Server is up on port 3000
```

When we run the script and the server is up and running (confirmed by the terminal message printed) we will notice that we are not brought back to the command prompt to run something else. This is because the node process is still up and running. With a web server this process is never going to stop unless we stop it, this is because its job is to stay up and running, listening and processing new incoming requests.

We can always shut the server down using control + c (or command + c) on our keyboard to shut down the web server process and bring us back to the command line prompt and we could always start it again if we need to.

When the server is up and running, we can visit it in the browser by following the URL link <http://localhost:3000/> and we will see the message Hello express! printed on the browser screen. When we visited that url in the browser, it went to our server and the Express server found the matching route for the root and it processed the request using our handler. The handler used res.send() to send back a text response and is exactly what we are seeing inside of the browser.

To add a second route, we would add a second call to app.get() but this time providing a different route string for the first argument. We can continue doing this to create even more routes to our web server.

Whenever we make changes to the server script file, we must restart the server for the changes to take effect because the server will not know that we have made changes to the file. This is simply shutting the server process using control + c on the keyboard and running the node command again. We can use something like nodemon package globally which will automatically restart our server whenever we make changes to our files.

```
:~/.../web-server$ nodemon src/app.js
```

If we installed nodemon locally to our project directory and not globally to our machine, we would have to run the following command to run the nodemon.js script locally passing in our file as an argument:

```
:~/web-server$ ./node_modules/nodemon/bin/nodemon.js src/app.js
```

We could setup a script within our package.json file so that we can run the npm run command followed by the script name to run the above within a shorter syntax.

If we were to navigate to a route that has not been setup within our express web server application, we would see a error message within our browser in the format of Cannot GET [routename] for example:

Cannot GET /about

We can setup a 404 page for routes that we do not have support setup for which we will learn in a later section as it is a more complex topic. We now know how to setup a simple Express web server and serve simple text to certain URL routes.

CHALLENGE:

Setup Two New Routes

1. Setup an about route and render a page title
2. Setup a weather route and render a page title

3. Test your work by visiting both in the browser

SOLUTION:

web-server > src > app.js:

```
1. const express = require('express')
2. const app = express( )
3. app.get('/about', (req, res) => {
4.     res.send('About Page')
5. })
6. app.get('/weather', (req, res) => {
7.     res.send('Weather Page')
8. })
9. app.listen(3000, () => {
10.    console.log('Server is up on port 3000')
11. })
```

```
:~/.../web-server$ node src/app.js
```

```
Server is up on port 3000
```

Visit <http://localhost:3000/about> and <http://localhost:3000/weather> to test your work.

Serving up HTML and JSON

We have explored how to send back a text string to the browser for a request; however, in reality we are never going to just send back a text string. We are either going to send back a HTML designed to be rendered in the browser or we are going to send back JSON data designed to be consumed and used by code.

To serve up HTML and JSON, we are going to once again use the `res.send()` method but we would change what we would pass inside as the argument.

HTML

We can add the html directly within the argument string as seen below using `<h1>` element tags:

```
1. app.get('/html', (req, res) => {  
2.   res.send('<h1>About Page</h1>')  
3. })
```

JSON

We add an either an object or an array as the value to the `.send()` method:

```
1. app.get('/object', (req, res) => {  
2.   res.send( { name: 'Elle', age: 20 } )  
3. })
```

```
1. app.get('/array', (req, res) => {  
2.   res.send( [ { name: 'Lisa' }, { name: 'Jason' } ] )  
3. })
```

Express is going to detect that we have provided either a object or an array as the send value and it will automatically go to stringify the JSON for us and it is going to get sent to the requestor in the JSON format correctly.

To conclude, we can use the .send() method to send back regular text, html and JSON data to the requestor using our web server. It is important to note in reality we will serve up the contents of an entire directory and therefore if we have a very large webpage with a lot of HTML, we would not want to write all of that code inside of a string inside our node server .send() method. We would rather have a separate HTML file which we will serve up instead.

Serving Up Static Assets

Typing HTML in a string in a JavaScript file is going to get out of hand really quickly as we add more to the given web page. We can write our HTML in separate .html files and have our express server to serve them. We can configure express to serve up an entire directory of assets which could contain HTML, CSS, JavaScript, videos, images, and more file types.

We can create a directory called public to store our application assets. Any file that goes in this directory will be served up as part of express server. We can create a index.html which is a special file and the name indicates that it is going to get served up by default and this is the page that is going to end up showing at the root of our website/web application.

web-server > public > index.html:

We need to teach express how to serve up the contents of a directory and one important information we need is the path to the public folder. This cannot be a relative path, it must be an absolute path from the root of our machine. To get this done, Node provides us with two handy variables called __dirname and __filename.

The __dirname is short for directory name returns the string path for the directory the current script lives in. This will provide an absolute path from the root of our hard drive all the way to the folder containing the script file. The __filename on the other hand returns a similar string to the __dirname; however, it provides path to the file itself (i.e. the directory and the file name). Both of these variables are provided by the wrapper main function.

We can use the __dirname variable to get the correct path to the public folder directory using path manipulation to get the absolute file path for the desired folder. There are two methods we can take, the first is using string manipulation and the second is to use the core Node module called Path that provides

us a tonne of great utilities for working with paths. The path module is cross platform compatible and is a useful tool to use compared to the string manipulation technique.

web-server > src > app.js:

1. const path = require('path')
2. const express = require('express')
3. const app = express()
4. app.use(express.static(path.join(__dirname, '../public')))

The path.join() is a function which allows us to return the final path. We pass to it the individual pieces to the path and it does the job of manipulating the string for us. We can use the __dirname to get a starting path and then add a second argument to manipulate that string. The double dots allows ('..') us to move up a directory from the __dirname path. We can move up multiple directories using the forward slash and double dots ('..../') like so which would move up by two directory and so on. To move into a folder we use the forward slash followed by the name of the folder we wish to move into. We can use the path module to easily string manipulate the directory of the __dirname from src directory to the public directory. We can use this to configure express to serve the public directory up.

We can use the .use() method to serve up the directory containing our static assets. The app.use() is a method to allow us to customise our server. We will explore the .use() method in more detail later. Within this method we need to call express.static() which is a function provided by express. Static takes in the path of the folder we want to serve up. This will setup our express server to serve the public directory.

If we now visit the root of our serve i.e. <http://localhost:3000/>, this will serve up our static file from our public directory i.e. the content from our index.html file which we setup in the public folder. We can also visit this html page by providing the explicit path i.e. <http://localhost:3000/index.html> which will provide the same result. It is important to note that index.html has a special meaning when related to web servers in which we can leave this path off of the URL if we are navigating to the root of the website. This ensures that something loads if the user does not explicitly provide a path.

This configuration also means that we are never going to see the route handler if we set one up for the root route i.e. app.get("", ...). This is because our express server is going to work through our application until it finds a match for that route. In the case of express.static() call, it is indeed going to find a match because it is going to find index.html within the public directory which is going to match the root URL because the file has a special name. This would mean the app.get() call for the root, will never actually run and therefore can be removed from the server script as it no longer serves much of a purpose.

If we have more HTML files within the public directory such as a about.html and help.html files, the express.static() function allows us to serve these files by visiting the route for these files without creating them with the .get() method. For example, we can visit the <http://localhost:3000/about.html> to serve the about.html static file and <http://localhost:3000/help.html> for the help.html file.

We have now explored the configuration setup of our Express server to serve static assets to our routes.

Serving Up CSS, JavaScript, Images and More

We have explored how to setup express to serve up contents from the public directory. In this section we are going to focus on how we can build out the public folder and adding other assets into the mix and using them in our HTML files.

The CSS file needs to be loaded by the browser which means it needs to go into the public folder as this is the only directory setup to be exposed by the web server. We can add our CSS files in the public folder, and we can stay organised by creating a CSS sub-directory and saving all of our CSS style sheets. This does not mean that the CSS will be loaded in the browser. We would need to load the CSS files inside of our HTML files first using the `<link rel="stylesheet" href="/css/style.css">` element within the `<head>` section. This will load in our CSS style sheet and style our HTML elements.

We can use either a relative (`href=".//css/style.css"`) or an absolute path (`href="/css/style.css"`) to link the css file. The absolute path starts with a forward slash / followed by the folder. The forward slash on our local machine would normally translate to our hard drive; however, the forward slash in this case is relative to the web server root which we setup as the public folder.

The same would apply for JavaScript, Images and more files. We can add these to the public directory within their own folders but must load them within our HTML files to view/serve them in the browser.

Dynamic Pages with Templating

We can use something called a template engine to render dynamic web pages using express. A template engine we will explore is called handlebars. Handlebars is going to allow us to two very important things. Firstly, it is going to allow us to render dynamic documents as opposed to static documents. Secondly, it is going to allow us to easily create code that we can reuse across pages such as a header and footer elements.

We will need to install two npm packages called handlebars (<https://www.npmjs.com/package/handlebars>) and hbs (<https://www.npmjs.com/package/hbs>).

Handlebars is a very low level library that implements handlebars in JavaScript. It can be used wide variety of settings such as the browser, server, desktop application via election or anywhere else that JavaScript can be used. To use handlebars in our express server, the module is not going to allow us to get this done; instead we need to install another handlebar library called hbs which is like a plugin for express which implements handlebars into express. Hbs uses handlebars behind the scenes and so we will only need to install hbs as it will install handlebars as its dependency.

```
:~/.../web-server$ npm i hbs@4.0.4
```

The process of having this setup is very easy. All we need to do is tell express which templating engine we installed and we do this by using a new method on our app express application variable.

web-server > src > app.js:

```
1. const path = require('path')
2. const express = require('express')
3. const app = express()
4. app.set('view engine', 'hbs')
```

The .set() method allows us to set a value for a given express setting and there are a few. We have a key which is the setting and the value which is the value we want to set for the setting. To setup the view we pass in the view engine as the first argument and the value we use as the second argument is the name of the module for our templating engine. This single line is all we need to get handlebars setup and we can actually use it to create dynamic templates.

When working with express, it expects all of our views, in this case the handlebars templates, to live in a specific folder. This is a specific folder at the root of the project called views.

web-server > views:

We would create all of our handlebars views within the views folder and we use the extension of .hbs to indicate a handlebars view file. For example we can create a index.hbs view file as seen below. A handlebar file is nothing more than HTML with a couple of nice to have little features for injecting dynamic values. This means we can start to use regular HTML syntax and create a regular HTML document to serve up.

web-server > views > index.hbs:

```
1. <!DOCTYPE html>
2. <html lang="en">
3.   <head>
4.     <script src="/js/app.js"></script>
5.     <title>Express Server - Serving Up Dynamic Templates</title>
6.   </head>
7.   <body>
8.     <h1>Weather</h1>
9.   </body>
10. </html>
```

So far the above code is static HTML code and is not taking advantage of handlebar's features; however, we can see how express can serve up the views template to our routes. Note that we can remove the index.html file from the public folder; however, it is not to say that we cannot serve static files. If we wish to serve static files, we would place them in the public folder, whereas if we need something more complex with dynamic content then handlebars (views folder) is our best option.

To serve up the hbs template, we need to setup a route, as seen below:

web-server > src > app.js:

```
1. ...
2. app.get('',(req, resp) => { res.render('index') } )
```

Inside of the anonymous callback function, we have so far explored sending the response back to the requestor via the `res.send()` method which would send back a string, HTML string, object or an array converted as a JSON data. To send a dynamic template to the requestor we use the `res.render()` method. The `.render()` method allows us to render one of our views which we configured express to use the hbs to use handlebars template engine to render our views. The first argument we need to pass to the `.render()` method is the name of the particular view we want to use and we do not need to add the file extension.

By calling `res.render()`, express goes off and gets the file which matches the name passed in as the argument to this method and it converts it into HTML and makes sure that the HTML gets back to the user in the browser. At this point we have a static hbs/HTML document and there is nothing dynamic about the file even though we are using handlebars. We are now going to explore handlebars features to see how we can make the file more dynamic.

The second argument to the `.render()` method is an object which contains all of the values we want that view to be able to access which we can use to make the template dynamic.

```
app.get('', (req, resp) => {
  res.render('index', {
    title: 'Weather App', name: 'John Doe'
  })
})
```

Within the view file we can use the handlebars syntax to make our templates dynamic and access the values accessible from the `.render()` method:

```
1. <!DOCTYPE html>
2. <html lang="en">
3.   <head>
4.     <script src="/js/app.js"></script>
5.     <title>Express Server - Serving Up Dynamic Template</title>
6.   </head>
7.   <body>
8.     <h1>{{ title }}</h1>
9.     <p>Created by {{ name }}</p>
10.    </body>
11. </html>
```

To inject values into a `.hbs` file, we open and close two curly brackets and inside of this we reference the value we are trying to access using the key name. In the above we are accessing the `title` key's value which is `Weather App`. `Weather App` which is the value will be injected into the `<h1>` element and displayed in the Browser.

We now know how to use handlebar templates to render some dynamic values that end up getting provided by Node.js.

Customising the Views Directory

We are going to explore how to customise how handlebars is setup so that we can use more advanced features of handlebars called partials to make our templates more dynamic. Currently, we have learned how to setup handlebars with our express server; however, currently the views folder must be in a specific location i.e. the root of the project and it must have the folder name of views.

We can change the folder name from views to something like templates; however, this will break our applications as express is not configured to look for a templates directory. The views folder and file path is the default location express expects the views directory to live in but we can customise this by telling express where to look. This will require us to create a brand new absolute path:

web-server > src > app.js:

1. ...
2. const viewsPath = path.join(__dirname, '../templates')
3. app.set('views', viewsPath)
4. app.set('view engine', 'hbs')

We would store the new absolute path in a variable and then tell express to use this path by using another app.set() method call. The setting we are configuring is called views and changing the default to our new viewsPath location. This will now enable us to use the templates folder as the new configured path. We now know how to customise the location and the name of the folder and we can change the location to anywhere we want and in a nested file structure as long as we update the viewsPath to point to it.

Everything that we have learned thus far about Express and its methods can be found on the Express website under the API reference page (<https://expressjs.com/en/4x/api.html>). This is a very useful API reference document which has 5 distinct categories. If we click on express() category we can see the methods we have available on express(). Under Applications we have all the methods we can use on the application object (which in all the above examples we have named app). This contains documentation on methods such as .get(), .listen(), .path(), .render(), .set(), etc. If we click on .set() we would see the parameters we can use with the .set() method and we should see views and view engine settings that we can set.

This is a very helpful document and we should always refer to this documentation to understand what parameters we can use for each methods we would want to use. Now that we understand how to use the documentation for Express, we should have a feel for how exactly the documentation works and how we can use it to discover new features and customise our own express application further.

Advanced Templating

Partials as the name suggests allows us to create a little template which is part of a bigger webpage. This is where partials come from. We can think of part of a webpage that we are going to reuse across multiple pages in our website. This will be things such as headers or footers where we want the exact same thing showing on every page to give our website a nice and unified feel. Partials is going to allow us to create, maintain and reuse markup throughout our website very easily without having to copy the markup between all the pages in our website.

To work with partials, we need to load in the hbs module and configure it:

web-server > src > app.js:

```
1. ...
2. const hbs = require('hbs')
3. const viewsPath = path.join(__dirname, '../templates/views')
4. const partialsPath = path.join(__dirname, '../templates/views')
5. app.set('views', viewsPath)
6. app.set('view engine', 'hbs')
7. hbs.registerPartials(partialsPath)
```

Once we load in the hbs module, we need to tell handlebars where we are going to put our partials. Partials are also files with .hbs extensions similar to the template views we have seen previously. We want to keep them separate so that we have one directory for our views and another for our partials. The structure we can adopt is a templates directory which has two sub-directories called views and partials.

To get things working we need to change the viewsPath to locate the template files within the views sub-directory. To tell handlebars where our partials are located we would create another variable to store the partials path location and then use the hbs.registerPartials() method to configure handlebar to look correct path for the partials directory. The configuration is now complete and we can now use partials to create more advanced templates.

Below is a simple partial example on how we can use partials with our view templates:

web-server > templates > partials > header.hbs:

1. <h1>{{ title }}</h1>

web-server > templates > partials > index.hbs:

1. {{>header}}

What goes inside of a partial is not a complete HTML document, instead it is just part of a webpage i.e. part of a mark up. To render a partial within our views file, we continue to use the two curly brackets like we did when we wanted to inject a value into the template in a specific location. For a partial, we need to add one additional character which is the greater than sign > followed by the partial name (i.e. the partial file name without the file extension or the file path) inside of the curly brackets.

If we are using nodemon to automatically restart our node web server whenever there is a change to files, we would need to make one slight change to our nodemon command, as currently this will not load in the new/updated .hbs files which will cause our application to crash. To fix this we would run the same command as before but adding the -e flag which stands for extensions. This allows us to provide a comma separated list of extensions that nodemon should watch. This will now monitor and restart the server for any changes to any .js and .hbs files we make in our project.

```
:~/.../web-server$ nodemon src/app.js -e js,hbs
```

We can continue to inject values into our partial files and this will continue to render dynamic text even when the partials are used inside of our rendered views templates. This allows us to break up our website into different smaller components which we can use across multiple pages of our website. This will allow us to write our code once which is easier to maintain and update in a single location which helps to scale our website should it grow in complexity.

404 Pages

If someone provides a URL that we have not explicitly setup in our Express server, we can show a 404 page instead of the generic Cannot GET /... text message coming from Express. We can use the 404 page to render a nice useful HTML along with a link back to the homepage so that the user can find a page that actually exists.

We need to setup another route handler using the app.get() method but place this right at the end just before starting up the server with the app.listen() method.

web-server > src > app.js:

1. ...
2. app.get('*', (req, res) => { app.render('404') })

In the first argument using the wildcard (*) means to match anything that has not been matched so far. This is why it is crucial to add this route towards the end of our code before we start the server up.

When Express gets an incoming request, it starts to look for a match starting from the top of our code working its way down in order. Therefore, it is going to look for the incoming route at the public directory first to see if there is any files matching the request to server and if it does not find any matching file it continues down the list looking at the next handler for a match. If it finds a matching handler with the incoming URL request, Express will render that handler's response and will not continue to look at the other routes below because it has found a matching route.

Where an incoming request does not match anything within the public directory and then all the routes defined, Express will finally come to the last defined route which contains the wildcard, which means everything is a match, and as a result this will return our 404 page back as a response to the requestor. This is all it takes to setup a route handler for a 404 page in Express.

We can take this further, and add another route combining the wildcard. This route must come before the final all capturing wildcard route to work:

```
app.get('/help/*', (req, res) => { app.send('Help article not found.') })
```

If we now follow the url link for example to <http://localhost:3000/help/test/> we will now see a more specific 404 message. We can therefore use the wildcard to match every request or match a bunch of requests that matching a specific pattern, in the above example matching anything after /help/ as long as it was not matched previously.